
Causal Validation and Behavioral Grounding of Moral Representation Directions in Language Models

Orion Reblitz-Richardson*

Abstract

Companion paper to [Reblitz-Richardson \(2026\)](#), which identified structured moral representations in language models via linear probing but acknowledged that probe directions are correlational. We provide three forms of validation that these directions are genuine features of the model’s computation. **Causal validation:** ablating a foundation’s direction from the residual stream at layer 12 specifically reduces log-probability of that foundation’s continuations (mean specificity -0.93 nats), while injecting directions at low amplitude ($\alpha = 1$) produces specific boosts that saturate into non-specific degradation at high amplitude — a dose-response curve consistent with real feature manipulation rather than noise injection. **Behavioral grounding:** projection-based 6-way foundation classification achieves 67.1% on held-out test pairs and 85.4% on causal evaluation prompts (chance = 16.7%), but only 26.7% on Moral Foundations Vignettes, revealing that real-world moral stimuli co-activate the care/harm direction regardless of their target foundation. **Sparse autoencoder analysis:** top-100 morally selective SAE features span a subspace that captures 19.6% of probe direction variance ($4.01\times$ random baseline, $p < 0.001$), indicating partial but significant convergence between unsupervised feature discovery and supervised direction extraction. Together, these results establish moral probe directions as causally load-bearing, behaviorally predictive features suitable for representation engineering interventions.

1 Introduction

Companion paper ([Reblitz-Richardson, 2026](#)) established that language models develop structured moral representations: six foundation-specific probe directions exhibit integration geometry (effective dimensionality 5, MFT-consistent clustering), and these directions are stable across extraction methods, dataset sizes, and text registers. But probe directions are correlational. A direction that separates moral from neutral text in activation space may reflect a confound (sentence length, formality, topic) rather than a genuine moral feature. Three questions remain open:

1. **Causal relevance.** Does the model *use* these directions during generation? If we ablate a foundation’s direction, does the model lose access to that specific moral concept? If we inject a direction, does the model shift toward that foundation’s content?
2. **Behavioral grounding.** Do the directions predict which moral foundation a novel text activates? Representational geometry tells us how the model *organizes* moral knowledge; behavioral benchmarking tests whether that organization is *functionally accessible* for moral reasoning.
3. **Mechanistic correspondence.** Do the supervised probe directions correspond to features the model discovers on its own? Sparse autoencoders (SAEs) learn unsupervised features from general text. If morally selective SAE features align with probe directions, the directions are not artifacts of the probing procedure but reflect native model structure.

*Distiller Labs. Correspondence to Distiller Labs <orion@orionr.com>.

We address all three questions using the same OLMo-2 1B model and 1,200-pair moral probing dataset from [Reblitz-Richardson \(2026\)](#). Section~4.1 presents direction ablation and steering injection experiments. Section~4.3 reports projection-based foundation classification on held-out, external, and causal evaluation stimuli. Section~4.4 compares SAE feature geometry with probe directions. Section~5 synthesizes the three forms of evidence and discusses implications for representation engineering of moral concepts.

2 Related Work

2.1 Causal methods in mechanistic interpretability

Activation patching ([Meng et al., 2022](#), [Vig et al., 2020](#)) and causal tracing ([Geiger et al., 2021](#)) establish whether identified features are functionally relevant by measuring behavioral changes under targeted interventions. Direction ablation (projecting out a direction from hidden states) tests necessity: does the model need this direction? Steering injection (adding a scaled direction to hidden states) tests sufficiency: can this direction control model behavior? [Arditi et al. \(2024\)](#) demonstrated that refusal behavior is mediated by a single direction, establishing the paradigm we apply to moral foundations. [Turner et al. \(2023\)](#) introduced activation addition for steering, showing that concept directions can modulate generation at inference time.

2.2 Representation engineering

[Zou et al. \(2023\)](#) proposed extracting concept directions via paired-difference PCA and using them for model control. [Marks and Tegmark \(2024\)](#) extended this to multi-dimensional concept spaces. Our work differs in validating directions extracted by three independent methods (probe weights, mean-difference, LEACE) and providing causal evidence that the directions are load-bearing for the specific moral concepts they encode.

2.3 Sparse autoencoders for feature discovery

SAEs decompose neural network activations into sparse, interpretable features ([Bricken et al., 2023](#), [Cunningham et al., 2023](#)). If supervised probe directions align with unsupervised SAE features, this provides evidence that the directions correspond to natural features of the model’s computation rather than artifacts of the probing procedure. Prior work has used SAEs primarily for safety-relevant features ([Templeton et al., 2024](#)); we apply them to structured moral concepts.

2.4 Moral reasoning in language models

[Reblitz-Richardson \(2026\)](#) established the geometric structure of moral representations that this paper validates causally. The Moral Foundations Vignettes ([Clifford et al., 2015](#)) provide an independent external stimulus set for behavioral benchmarking.

3 Methods

3.1 Model and dataset

All experiments use OLMo-2-0425-1B (16 layers, 2048 hidden dim) with mean-difference directions extracted from a 1,200-pair moral probing dataset (200 pairs per MFT foundation, 960 train / 240 test, three text registers: declarative, situational, narrative). Directions are normalized to unit length. See [Reblitz-Richardson \(2026\)](#) for dataset construction and direction extraction.

3.2 Causal validation methods

3.2.1 Direction ablation

For each foundation f and layer ℓ , we register a forward hook that projects out the foundation direction $\mathbf{d}_f^{(\ell)}$ from the hidden state:

$$\mathbf{h} \leftarrow \mathbf{h} - (\mathbf{h} \cdot \mathbf{d}_f^{(\ell)}) \mathbf{d}_f^{(\ell)}$$

We then measure the change in log-probability of target continuations from a 48-prompt causal evaluation set (8 prompts per foundation, with labeled target and off-target continuations).

Metrics. For each (ablated foundation, layer):

- *On-target* Δ : mean log-prob change on prompts whose target foundation matches the ablated direction.
- *Off-target* Δ : mean log-prob change on prompts from other foundations.
- *Specificity*: on-target Δ minus off-target Δ . Negative specificity indicates that ablation specifically harms the target foundation.

3.2.2 Steering injection

For each foundation f , layer ℓ , and amplitude $\alpha \in \{1, 2, 5, 10, 20\}$, we add a scaled direction to the hidden state:

$$\mathbf{h} \leftarrow \mathbf{h} + \alpha \mathbf{d}_f^{(\ell)}$$

and measure log-probability changes on the same evaluation set. The dose–response curve across α values distinguishes genuine feature manipulation (specificity at low α , saturation at high α) from noise injection (monotonic degradation).

3.2.3 Causal evaluation prompts

The 48-prompt evaluation set comprises three formats: 24 completion prompts (4 per foundation) with a sentence stem and 3–4 candidate continuations, 12 forced-choice prompts (2 per foundation), and 12 natural prompts (2 per foundation). Each prompt has labeled target and off-target continuations with foundation labels, enabling measurement of foundation-specific log-probability changes under ablation and injection.

3.3 Behavioral grounding methods

3.3.1 Projection-based classification

For each input text, we collect the last-token residual stream activation at target layers (4, 8, 12), project onto all six foundation directions, and average across layers. The foundation with the highest projection is the predicted label. We also test a *debiased* variant that subtracts the mean projection across all six foundations before classification, removing the shared moral-salience component.

3.3.2 Evaluation sets

1. **Held-out test set** (240 pairs, 40 per foundation): internal validation from the probing dataset.
2. **Moral Foundations Vignettes** (30 items, 5 per foundation): curated from [Clifford et al. \(2015\)](#), providing external validation with established MFT stimuli independent of our dataset pipeline.
3. **Causal evaluation prompts** (48 prompts): cross-validation with the causal experiments, testing whether directions that are causally relevant also predict foundation identity.

3.4 Sparse autoencoder analysis methods

3.4.1 SAE training

We train a ReLU sparse autoencoder (16,384 latent dimensions, $8\times$ expansion) on 2M token activations from the C4 corpus ([Raffel et al., 2020](#)) at layer 8, using L1 sparsity regularization ($\lambda = 0.005$). The decoder columns are constrained to unit norm. The pre-encoder bias is initialized to the mean activation.

3.4.2 Moral feature identification

We encode 960 moral and 960 neutral sentences through the trained SAE and compute per-feature selectivity: the mean activation difference between moral and neutral inputs. Top- k features (ranked by $|\text{selectivity}|$) are compared with probe directions via:

- **Individual alignment:** cosine similarity between each feature’s decoder column and each foundation direction.
- **Subspace overlap:** fraction of probe direction variance captured by the top- k SAE feature subspace (via SVD projection).

3.4.3 3.4.3 Random baseline

To establish significance of subspace overlap, we compute a null distribution: 1,000 iterations of projecting probe directions onto 100 random unit vectors in $\mathbb{R}^{2048}$ and recording mean membership.

4 Results

4.1 Direction ablation is foundation-specific

Ablating a foundation’s direction from the residual stream specifically reduces log-probability of that foundation’s continuations while leaving other foundations largely unaffected (Table~??). Specificity (on-target Δ minus off-target Δ) is negative for all 18 (foundation $\times$ layer) combinations, confirming that each direction carries information specific to its foundation.

Layer dependence. Specificity increases with depth. At layer 4, mean specificity across foundations is -0.30 ; at layer 8 it is -0.43 ; at layer 12 it is -0.93 . The strongest effects are sanctity at layer 12 (specificity = -1.89 , on-target $\Delta = -2.03$ nats) and loyalty at layer 12 (-1.08). This pattern—ablation most damaging in later layers—is consistent with the model relying more heavily on moral representations as they approach the output.

Foundation heterogeneity. Sanctity and loyalty directions are the most causally load-bearing (largest specificity at layer 12), while fairness and authority show moderate effects. Care shows strong on-target effects but also larger off-target perturbation, consistent with care/harm co-activating with other moral content (the “care saturation” phenomenon observed in behavioral benchmarking, §4.3).

4.2 Steering injection shows dose–response specificity

At low amplitude ($\alpha = 1$), injecting a foundation direction into the residual stream at layer 8 produces a *positive* on-target boost (mean $+0.20$ nats across foundations) with near-zero off-target effect (mean $+0.03$), yielding positive specificity for all six foundations (Table~??).

Dose–response curve. As α increases:

- $\alpha = 1$ –2: specific boost. On-target log-probability increases, off-target is near baseline. This is the “therapeutic window.”
- $\alpha = 5$: mixed. Some foundations maintain positive on-target (sanctity $+0.25$, loyalty $+0.07$), others turn negative. Specificity remains positive for all.
- $\alpha = 10$ –20: non-specific degradation. Both on-target and off-target log-probabilities decrease, but off-target decreases *more*, so specificity remains positive ($+2.4$ at $\alpha = 10$, layer 8). The model is pushed out of its functional range.

This dose–response pattern distinguishes the moral directions from random directions: noise injection would produce monotonic degradation at all amplitudes, not the low- α specific boost we observe.

Layer dependence. Layer 12 maintains specificity to higher α values than layers 4 and 8. At $\alpha = 5$, all six foundations show positive on-target Δ at layer 12 (mean $+0.25$), while layers 4 and 8 are already showing degradation for some foundations. This robustness at layer 12 parallels the strong ablation effects, confirming that later layers are the primary locus of moral direction utilization.

4.3 Behavioral grounding: projection predicts foundation identity

Projection-based 6-way classification (debiased) achieves:

Dataset	Accuracy	Chance
Causal prompts	85.4% (41/48)	16.7%
Held-out test set	67.1% (161/240)	16.7%
Moral Foundations Vignettes	26.7% (8/30)	16.7%

4.3.1 4.3.1 Test set performance

Per-foundation accuracy is uneven: care (87.5%), sanctity (75.0%), loyalty and authority (72.5%) perform well, while fairness (45.0%) and liberty (50.0%) are weaker. The confusion matrix reveals that fairness and liberty errors scatter across multiple foundations rather than concentrating on a single confound, suggesting these foundations have more distributed representations that share variance with multiple other directions.

4.3.2 4.3.2 External validation: Moral Foundations Vignettes

The MFV accuracy of 26.7% (above the 16.7% chance baseline) masks a systematic pattern: the confusion matrix shows that care/harm dominates classification regardless of the true foundation. All five care items are correctly classified, but 23 of 25 non-care items are also classified as care.

This is not a failure of the directions but a property of the stimuli. The MFV items are descriptions of witnessing morally transgressive acts (You see a boy kicking a puppy'), and witnessing harm *genuinely activates* the care/harm representation regardless of which foundation the scenario is designed to target. A fairness violation (You see a judge giving a lighter sentence to a wealthy defendant') also involves potential harm to the disadvantaged party. An authority violation ("You see a student cursing at a teacher') also activates interpersonal harm sensitivity.

This co-activation is consistent with the integration geometry: the shared moral-salience component (mean pairwise cosine ~ 0.3 between foundation directions) means that stimuli activating *any* foundation will partially activate *all* foundations, with care receiving disproportionate activation due to the harm content present in most moral scenarios. Debiasing (subtracting the mean projection) does not resolve this because the care co-activation is genuine, not an artifact of the shared component.

4.3.3 4.3.3 Causal prompt performance

The 85.4% accuracy on causal evaluation prompts—which were designed for targeted foundation activation, not harm witnessing—confirms that the directions are behaviorally predictive when stimuli are foundation-specific. Per-foundation accuracy is uniformly high (75–87.5%), with no care dominance.

4.4 SAE features partially recover moral subspace

4.4.1 4.4.1 Training results

The layer-8 SAE achieves $L0 = 1,932$ (11.8% of 16,384 features active per token) and $FVU = 0.285$ (71.5% variance explained) after 3 epochs on 2M tokens. The moderate FVU reflects limited training scale; production SAEs typically train on 100M+ tokens.

4.4.2 4.4.2 Moral selectivity

Only 1 of 16,384 features has moral selectivity $|s| > 0.1$ (mean activation on moral minus neutral text). Moral information is distributed across many features rather than concentrated in a few—consistent with the high effective dimensionality of moral representations found in [Reblitz-Richardson \(2026\)](#).

4.4.3 4.4.3 Subspace overlap

The top-100 morally selective SAE features, despite individually showing no alignment with probe directions (0% of features with $|\cos| > 0.2$), collectively span a subspace that captures 19.6% of mean-difference direction variance. The random baseline is $4.89\% \pm 0.36\%$ (matching the analytical expectation of $100/2048 = 4.88\%$), yielding a ratio of $4.01 \times$ random ($p < 0.001$ by permutation).

Probe-weight directions show weaker overlap (9.1%, $\sim 2\times$ random), consistent with LEACE and probe-weight directions capturing partially different aspects of the moral subspace.

This result has two interpretations. First, even a partially trained SAE discovers features that overlap with supervised moral directions, suggesting these directions reflect genuine structure in the model’s representations rather than probing artifacts. Second, the overlap is modest (19.6%), indicating that either (a) the SAE needs substantially more training to decompose moral features cleanly, or (b) moral representations are inherently distributed across many low-selectivity features that resist separation by current SAE methods.

5 Discussion

5.1 Three converging lines of evidence

The causal, behavioral, and mechanistic results converge on a single conclusion: the moral probe directions identified in [Reblitz-Richardson \(2026\)](#) are genuine features of the model’s computation, not artifacts of the probing procedure.

Direction ablation establishes *necessity*: the model needs these directions to generate foundation-specific content. Steering injection establishes *sufficiency*: adding these directions shifts model behavior toward the target foundation, with a dose–response curve characteristic of real feature manipulation. Behavioral grounding establishes *functional accessibility*: the directions predict which foundation a novel text activates, at accuracy levels far above chance. SAE analysis establishes *mechanistic correspondence*: unsupervised feature discovery partially recovers the same subspace that supervised probing identifies.

No single line of evidence would be conclusive. Ablation could reflect a generic disruption effect; injection could be noise-driven; behavioral accuracy could reflect confounds; SAE overlap could be coincidental. But taken together, the convergence across independent methods with different assumptions and failure modes provides strong evidence for the reality of the moral direction structure.

5.2 The care saturation phenomenon

The MFV results reveal a property of real-world moral stimuli that laboratory-controlled probing datasets obscure: witnessing moral transgressions of *any* kind activates the care/harm representation. This is not a deficiency of the probing directions but a genuine feature of moral cognition as encoded by the model. Moral psychology has long recognized that harm is the most accessible moral dimension cross-culturally ([Graham et al., 2013](#)), and the model’s representations reflect this asymmetry.

This has practical implications for steering applications: injecting a non-care foundation direction into a context that already contains harm content will contend with strong care co-activation. Effective foundation-specific steering may require *suppressing* the care component while boosting the target foundation—a two-direction intervention rather than a single injection.

5.3 Layer-dependent causal roles

Ablation and injection reveal complementary layer-dependent patterns:

- **Ablation specificity increases with depth** (strongest at layer 12), indicating that later layers rely more heavily on foundation-specific information for generation.
- **Injection at layer 12 tolerates higher α** before saturating, consistent with later-layer representations being more robust to perturbation.

These patterns suggest a processing hierarchy: earlier layers encode moral content in a form that is more malleable but less causally connected to output, while later layers encode it in a form that is more rigid but directly drives generation.

5.4 Toward a steering fitness function

The results suggest a concrete fitness function for evaluating moral directions as steering targets:

1. **Ablation specificity** < -0.5 at the target layer (direction is causally load-bearing).
2. **Injection specificity** > 0 at $\alpha = 1$ (direction can produce specific behavioral shifts).
3. **Behavioral accuracy** $> 50\%$ on held-out data (direction is functionally discriminative).
4. **SAE subspace overlap** $> 2\times$ random baseline (direction corresponds to native features).

Sanctity, loyalty, and care directions pass all four criteria. Fairness and liberty pass criteria 1–3 but show weaker SAE overlap, suggesting their representations may be more distributed. Authority passes all criteria at moderate levels.

5.5 Limitations

SAE training scale. The 2M-token SAE achieves only 71.5% variance explained. Production-scale SAE training (100M+ tokens) would likely increase both feature quality and moral direction overlap. The current results should be interpreted as a lower bound.

Single model. All experiments use OLMo-2 1B. Generalization to larger models and different architectures is untested.

Causal evaluation set size. The 48-prompt evaluation set (8 per foundation) limits statistical power for per-foundation claims. The direction of all effects is consistent, but effect sizes should be interpreted with appropriate uncertainty.

Steering injection confounds. At high α , injection degrades model performance broadly. The positive specificity at high α reflects *relative* preservation of on-target content, not *absolute* improvement. The therapeutically relevant regime is $\alpha = 1-2$.

6 Conclusion

We have provided three independent forms of evidence that the moral probe directions identified in Reblitz-Richardson (2026) are genuine features of OLMo-2 1B’s computation. Direction ablation establishes necessity (specificity -0.93 at layer 12); steering injection establishes sufficiency (specific boost at $\alpha = 1-2$ with dose–response saturation); and sparse autoencoder analysis establishes mechanistic correspondence ($4.01\times$ random subspace overlap). Behavioral benchmarking confirms that these directions are functionally accessible, achieving 67.1% accuracy on held-out test data and 85.4% on causal evaluation prompts.

The care saturation phenomenon—care/harm activation dominating classification of real-world moral stimuli regardless of target foundation—is not a limitation of the directions but a genuine property of moral representation that has implications for steering applications. Effective moral steering may require multi-direction interventions that suppress care co-activation while boosting the target foundation.

These results transform the moral geometry from a descriptive finding into a tool for representation engineering. The directions are not just *readable* from the model; they are *writable* into the model with predictable behavioral effects. Future work should test generalization to larger models, refine the SAE analysis with production-scale training, and develop the multi-direction steering strategies that the care saturation phenomenon motivates.

Acknowledgments and Disclosure of Funding

This work made extensive use of Anthropic’s Claude (the Claude Code agent on Opus 4.6) for code scaffolding, experimental scripts, and prose drafting. The author retains responsibility for experimental design, all scientific claims, and final wording.

References

Andy Arditi, Oscar Obeso, Aaquib Syed, Daniel Paleka, Nina Panickssery, Wes Gurnee, and Neel Nanda. Refusal in language models is mediated by a single direction. *arXiv preprint arXiv:2406.11717*, 2024. URL <https://arxiv.org/abs/2406.11717>.

- Trenton Bricken, Adly Templeton, Joshua Batson, Brian Chen, Adam Jermyn, Tom Conerly, Nick Turner, Cem Anil, Carson Denison, Amanda Askeel, Robert Lasenby, Yifan Wu, Shauna Kravec, Nicholas Schiefer, Tim Maxwell, Nicholas Joseph, Zac Hatfield-Dodds, Alex Tamkin, Karina Nguyen, Brayden McLean, Josiah E. Burke, Tristan Hume, Shan Carter, Tom Henighan, and Christopher Olah. Towards monosemanticity: Decomposing language models with dictionary learning. *Transformer Circuits Thread*, 2023. URL <https://transformer-circuits.pub/2023/monosemantic-features>.
- Scott Clifford, Vijeth Iyengar, Roberto Cabeza, and Walter Sinnott-Armstrong. Moral foundations vignettes: A standardized stimulus database of scenarios based on moral foundations theory. *Behavior Research Methods*, 47(4):1178–1198, 2015. doi: 10.3758/s13428-014-0551-2.
- Hoagy Cunningham, Aidan Ewart, Logan Riggs, Robert Huben, and Lee Sharkey. Sparse autoencoders find highly interpretable features in language models. *arXiv preprint arXiv:2309.08600*, 2023. URL <https://arxiv.org/abs/2309.08600>.
- Atticus Geiger, Hanson Lu, Thomas Icard, and Christopher Potts. Causal abstractions of neural networks. *Advances in Neural Information Processing Systems*, 34, 2021.
- Jesse Graham, Jonathan Haidt, Sena Koleva, Matt Motyl, Ravi Iyer, Sean P. Wojcik, and Peter H. Ditto. Moral foundations theory: The pragmatic validity of moral pluralism. In *Advances in Experimental Social Psychology*, volume 47, pages 55–130. Academic Press, 2013.
- Samuel Marks and Max Tegmark. The geometry of truth: Emergent linear structure in large language model representations of true/false datasets. *arXiv preprint arXiv:2310.06824*, 2024. URL <https://arxiv.org/abs/2310.06824>.
- Kevin Meng, David Bau, Alex Andonian, and Yonatan Belinkov. Locating and editing factual associations in GPT. *Advances in Neural Information Processing Systems*, 35, 2022.
- Colin Raffel, Noam Shazeer, Adam Roberts, Katherine Lee, Sharan Narang, Michael Matena, Yanqi Zhou, Wei Li, and Peter J. Liu. Exploring the limits of transfer learning with a unified text-to-text transformer. *Journal of Machine Learning Research*, 21(140):1–67, 2020.
- Orion Reblitz-Richardson. The moral geometry of language models: Integration structure in foundation-specific representations. *arXiv preprint*, 2026. Companion paper.
- Adly Templeton, Tom Conerly, Jonathan Marcus, Jack Lindsey, Trenton Bricken, Brian Chen, Adam Pearce, Craig Citro, Emmanuel Ameisen, Andy Jones, Hoagy Cunningham, Nicholas L. Turner, Callum McDougall, Monte MacDiarmid, C. Daniel Freeman, Theodore R. Sumers, Edward Rees, Joshua Batson, Adam Jermyn, Shan Carter, Chris Olah, and Tom Henighan. Scaling monosemanticity: Extracting interpretable features from Claude 3 Sonnet. *Transformer Circuits Thread*, 2024. URL <https://transformer-circuits.pub/2024/scaling-monosemanticity>.
- Alexander Matt Turner, Lisa Thiergart, David Udell, Gavin Leech, Ulisse Mini, and Monte MacDiarmid. Activation addition: Steering language models without optimization. *arXiv preprint arXiv:2308.10248*, 2023. URL <https://arxiv.org/abs/2308.10248>.
- Jesse Vig, Sebastian Gehrmann, Yonatan Belinkov, Sharon Qian, Daniel Nevo, Yaron Singer, and Stuart Shieber. Investigating gender bias in language models using causal mediation analysis. *Advances in Neural Information Processing Systems*, 33, 2020.
- Andy Zou, Long Phan, Sarah Chen, James Campbell, Phillip Guo, Richard Ren, Alexander Pan, Xuwang Yin, Mantas Mazeika, Ann-Kathrin Dombrowski, Shashwat Goel, Nathaniel Li, Michael J. Byun, Zifan Wang, Alex Mallen, Steven Basart, Sanmi Koyejo, Dawn Song, Matt Fredrikson, J. Zico Kolter, and Dan Hendrycks. Representation engineering: A top-down approach to AI transparency. *arXiv preprint arXiv:2310.01405*, 2023. URL <https://arxiv.org/abs/2310.01405>.

Appendices

Supplementary material.

A Full Ablation Tables

Tables ??–?? report the complete ablation results for all six foundations at layers 4, 8, and 12. Each cell shows the mean change in log-probability (nats) under direction ablation, with on-target and off-target columns.

[Full tables to be generated from experimental output JSON.]

B Steering Dose–Response Curves

Figures ??–?? show the full dose–response curves for all six foundations at layers 4, 8, and 12. Each panel plots on-target and off-target log-probability change as a function of injection amplitude $\alpha \in \{1, 2, 5, 10, 20\}$.

[Dose–response figures to be generated from experimental output JSON.]

C Behavioral Confusion Matrices

Figure ?? shows the confusion matrix for projection-based 6-way classification on the held-out test set (240 pairs). Figure ?? shows the corresponding matrix for the Moral Foundations Vignettes (30 items), illustrating the care saturation phenomenon: 23 of 25 non-care items are classified as care.

[Confusion matrices to be generated from experimental output JSON.]

D SAE Feature Details

D.1 Training hyperparameters

Parameter	Value
SAE width	16,384
Expansion factor	$8\times$
Hidden dim	2,048
Training tokens	2M (C4)
Batch size	4,096
Learning rate	3×10^{-4}
L1 coefficient	5×10^{-3}
Epochs	3
L0 (final)	1,932
FVU (final)	0.285

Table 1: SAE training configuration and final metrics.

D.2 Per-foundation subspace overlap

[Per-foundation overlap table to be generated from experimental output JSON.]

D.3 Random baseline distribution

The null distribution (1,000 iterations of projecting probe directions onto 100 random unit vectors in $\mathbb{R}^{2048}$) has mean 0.0489 ± 0.0036 , matching the analytical expectation of $100/2048 = 4.88\%$. The observed overlap of 19.6% corresponds to $p < 0.001$ (no random iteration exceeded the observed value).