


Comprehension (harmfulness) and refusal are separately encoded in reasoning models

Harmfulness is strongly encoded — peaks at the instruction token



...yet nearly orthogonal to refusal — separate concepts

