
Knowing Without Doing

A Pre-Registered Panel of the Gap Between a Language Model’s Stated Moral Judgment and Its Action

Orion Reblitz-Richardson*

Abstract

A model can state that an action is wrong and take it anyway. We measure that gap directly, scenario by scenario, on an open 7B chat model, with the model’s own stated judgment as the reference and a discrete, logged action as the outcome. The panel is pre-registered: 248 scenarios across five pressure families (task completion, sycophancy under pushback, prohibited shortcuts, in-group favor, third-party harm), each written in a matched agent frame and a third-person judgment frame, each with a pressure-removed twin and three paraphrased judgment frames, half written by each of two generator models and cross-labeled by the other.

On OLMo-3-7B-Instruct the model takes the action it judged wrong on one screened scenario in five by majority vote (0.19, 95% CI 0.13 to 0.28) and on two rollouts in five. The measurement sits inside a calibration ladder: it exceeds the same statistic on pressure-removed twins by 0.10 (0.02 to 0.18), and a known-gap control in which the system prompt orders the violating action reaches 0.58. The gap survives a paraphrase-majority judgment reference. On the majority readout it is not resolved against the strictest reference, where all four judgment frames must agree; on a pre-registered secondary readout that uses the saved next-token probabilities instead of majority votes, the excess over the null holds at every strictness level and is largest at the strictest (0.08, 0.03 to 0.13).

The gap has no family structure at this power: third-party harm is not the family where it is smallest, so the action channel is not organized by the harm content that refusal reads. The same gap appears in the base model’s raw completion frame at least as strongly as in the instruct model’s, on the same scenarios, so post-training does not install it. One family depends on which model wrote its scenarios; a blind human read finds no construction asymmetry, and the difference is graded rather than reversed. We release the panel, the harness, the pre-registration with its sixteen dated amendments, and the per-rollout arrays.

1 Introduction

A language model can state, when asked as an observer, that an action is wrong, and then take that action when placed inside the situation as the agent. Whether that happens, how often, and whether it depends on what the model comprehends about the situation are questions a mechanistic program about moral representations needs answered before it can ask what the model’s *action* reads from its representations. Our previous work established what the *refusal* decision reads on the same models: a low-rank harm slice of a broader moral subspace, written into a narrow control-token channel (Reblitz-Richardson, 2026). This paper builds the behavioral target for the next step, action selection, and measures it.

*Distiller Labs. Correspondence to Distiller Labs <orion@orionr.com>.

The quantity is a knowing–doing gap: per scenario, the divergence between what the model judges to be right, elicited in a third-person frame, and what it does, elicited in a matched agent frame with a discrete, logged action. The reference for “right” is the model’s own stated judgment, not an external label, so the construct is self-inconsistency and sidesteps the contestability of external ethics. That construct is not new. Huang et al. (Huang et al., 2026) and Shen et al. (Shen et al., 2025) measure gaps between a model’s stated values and its enacted choices; Rakshit et al. (Rakshit et al., 2026) add a fast-versus-slow deliberation contrast; Gu et al. (Gu et al., 2025) compare stated and revealed preferences; Hosseini et al. (Hosseini et al., 2026) find a judgment–consequence gap in clinical allocation; and in non-moral domains Cheng et al. (Cheng et al., 2026) and Basu et al. (Basu et al., 2026) use a knowing–doing gap as the target that mechanistic interventions are scored against. What this paper adds is narrower and, we think, more useful for that last purpose: a per-scenario, self-referenced *moral* gap measured with the model as the agent under typed pressure families, inside a calibration ladder that bounds how much of any gap is reference noise or frame change, with a harm-involving family that ties back to what refusal reads, and a base-versus-instruct comparison in one raw-completion frame that none of the prior panels include.

The panel was pre-registered before any scenario existed, and every construction decision made afterwards is a dated amendment in the same document (Section A); there are sixteen. We report three pods on one model, OLMo-3-7B-Instruct with its base checkpoint (Team OLMo, 2025): a 96-scenario pilot that tested the instrument, a 320-scenario full panel, and a 120-scenario second round that lifted the screened count past the pre-registered gate and ran a paraphrase-swap cell on the one family whose result depended on which model had written it. Scenarios were generated half by Claude and half by GPT, cross-labeled by the other; the harness that parses actions was calibrated twice on real replies against two independent judges.

The findings, in the order the ladder licenses them. On the screened panel the model takes the action it judged wrong on one scenario in five by majority vote (0.19, 95% CI 0.13 to 0.28) and two rollouts in five; the excess over the same statistic on pressure-removed twins is 0.10 (0.02 to 0.18), and a control in which the system prompt orders the violating action reaches 0.58. The gap survives a paraphrase-majority judgment reference; on the majority readout it is not resolved against the strictest reference, where all four judgment frames must agree, and on a pre-registered secondary readout built from the saved next-token probabilities it holds at every strictness level and is largest at the strictest (0.08, 0.03 to 0.13). The gap has no family structure at this power: third-party harm is not where it is smallest, so the action channel is not organized by the harm content refusal reads. The same gap is present in the base model’s raw completion frame at least as strongly as in the instruct model’s on the same scenarios, so post-training did not install it. One family’s rate depends on which model wrote its scenarios; a blind human read finds no construction asymmetry, and on the continuous readout the difference is graded, not reversed.

The paper is organized around the instruments rather than the story. Section 2 describes the panel; Section 3 the harness, the ladder, and the two readouts; Section 4 the gap and its robustness to the judgment reference; Section 5 the families and the generator effect; Section 6 the base-versus-instruct cells; Section 7 what the result does and does not say about the action channel.

2 The panel

2.1 Scenarios and frames

A scenario is a self-contained situation with a moral norm at stake, a situational pressure that makes violating it locally advantageous, and two to four discrete options, at least one norm-consistent and one norm-violating; most scenarios also carry a neutral option (hold, ask, escalate). Five families vary the pressure type while holding the norm class distribution roughly constant (Table 1); they draw on the pressure genres of the MACHIAVELLI benchmark (Pan et al., 2023), the agentic-misalignment scenarios (Lynch et al., 2025), and sycophancy-under-pushback evaluations (Sharma et al., 2023), rewritten as short closed-menu situations. Every scenario is written twice: an **agent frame** in the second person (“You are Dana, the release coordinator. . .”) for the action readout, and a **third-person frame** about the same person for the judgment readout, with the same entities, facts, sentence order, and length within 15% (tokenized with the model’s own tokenizer); only pronouns and the role framing differ. Option texts are plain descriptions of actions with no evaluative words, presented under neutral letters whose order is a fresh seeded permutation on every rollout.

Table 1: The five pressure families. F2 is an appendix family by pre-registration: its action is closer to a second judgment than to doing something, and it never entered the gate.

id	family	pressure	action surface
F1	task completion	goal pursuit vs honesty (report a failure or fabricate success)	option choice
F2	social cost	truth vs approval (user pushes back without new evidence)	multi-turn; final answer is the action
F3	instrumental	rule vs shortcut (a prohibited but efficient route)	tool selection from a fixed menu
F4	loyalty/fairness	in-group favor vs fair treatment	allocation choice
F5	third-party harm	harm to a non-present party vs task success	option choice

Each scenario also carries a **pressure-removed twin** (both frames with the incentive sentence replaced by mundane filler, options unchanged) for the matched-null rung, a **harm-matched twin** for F1 through F4 (same harm level, a different non-harm norm at stake) so the harm-family contrast is not a harm-level contrast, and, from the full panel on, **three paraphrases** of each third-person frame for the reference-strictness ladder. Covariate tags record harm level (rated by a non-generator model), norm class, register, the presence of a named user and of a tool menu, and the generator.

2.2 Generation and labeling

Half of every family was written by Claude (Opus 5) and half by GPT (5.5), under one fixed system prompt whose hash is stored with every scenario; a family-level result that differs by generator is an anomaly by pre-registered rule, not a finding, and is reported per generator before it is pooled. Each scenario’s **external label** (which option an outside rater calls right, and a harm level) was produced by the *other* provider, with the script refusing same-provider pairs. Drafts were validated against the construction rules and regenerated with the violations fed back; safety refusals by a generator were regenerated in-plan or, for four paraphrase requests, handed to the other generator with the exception recorded on the scenario.

The panel of record has 248 primaries (40 per family plus 16 added to F1, F3, and F5 in a second round), 152 harm twins, and 40 swapped F4 scenarios used only by one cell (Section 5). Three GPT-written F4 twins whose option labels the cross-rater found inverted are flagged and excluded from every analysis.

2.3 Readouts

Judgment $J(s)$: the third-person frame, one shot (“Which option should Dana choose, and why?”), parsed to an option; eight rollouts at temperature 0.7 plus one greedy. $J(s)$ is the greedy option unless the sampled majority disagrees, in which case the scenario is flagged unstable and excluded, counted. **Action** $D(s)$: the agent frame, chat template, the model as the actor, “reply with only the letter”; 32 rollouts at temperature 0.7, reported as the violating fraction and, for the binary gap, by majority. The decision position is the first generated token, and the full next-token distribution there is saved for every rollout. **Raw frame**: the scenario, the option list, and a fixed prefix ending before the option token; the readout is the normalized log-probability over the option letters, taken on both the instruct and the base model, with a mass floor (0.5 on the option letters) below which the model is not engaging the option set.

The gap. $KDG(s) = 1$ if the majority action is norm-violating while the stable judgment named a non-violating option; 0 if the majority action matches the judgment; undefined, counted, otherwise. A neutral option counts as non-violating for the reference (amendment A11). The rate is the mean over screened scenarios with a nested bootstrap over scenarios and rollouts.

3 Instruments: harness, screen, ladder, and two readouts

3.1 Harness and calibration

One parser and one classifier (kdg_harness 1.0.0) serve every cell, so no cell can drift onto a private classifier. The parser resolves a reply to an option by a fixed ladder of rules (answer line, bare letter, tool name, leading letter, unique mention, contrastive commitment) and leaves refusals, hedges, and unmarked two-option mentions unparsed rather than guessed. It was calibrated in two stages: on 200 synthetic reply formats built on the panel’s own scenarios, where it agreed with the construction labels on all 200 and with an independent judge on 197; and, before any verdict, on 200 real replies from each pod, oversampling every reply the parser had not resolved by a clean rule, against two independent judges (one Claude, one GPT). Agreement was 0.99 and 1.00 with the judges and 0.99 between them on the full panel, with the harness the conservative side of every disagreement (Section B). On the pods of record the parse rate was 0.9997 on 10,240 action rollouts and 1.000 on 2,880 judgment replies.

3.2 Screen and gates

The panel is useful only where pressure engages. A scenario passes the **screen** if its judgment is stable and its violating fraction across the 32 action rollouts lies in $[0.15, 0.85]$ (mixed outcomes), or exceeds 0.85 with a stable norm-consistent judgment (a clean gap); scenarios the model never violates (no pressure) or never judges stably (no reference) are dropped and counted. The pre-registered **pilot gate** required at least 14 of 48 gate-family primaries to pass with at least two families contributing; the **full gate** requires at least 60 screened primaries across at least three gate families, a family whose gap CI excludes zero, harness agreement of at least 0.95, and no family whose result reverses across generator. The screen’s binomial properties at 32 rollouts (a true 0.10 scenario lands in the mixed band one time in five) are why the panel’s unit of inference is the scenario-level bootstrap, not the screen label.

3.3 The calibration ladder

No gap or absence of a gap is stated bare. Every headline sits inside four rungs measured on the same instrument and model. The **floor** is re-elicitation noise on the reference: the greedy judgment against the greedy judgment on a paraphrased frame. The **matched null** is the gap statistic on the pressure-removed twins: what frame change alone, with no incentive, produces. The **measurement** is the statistic on the screened panel. The **positive band** is the statistic under a known-gap control in which the system prompt orders the violating action: what the instrument returns when a gap is forced. Verdict sentences carry their rung and their detection bar; comparisons are paired difference CIs on the same scenarios, never CI overlap.

3.4 Reference strictness

A self-referenced gap is only as good as its reference. The judgment reference is therefore elicited on four third-person frames (the original and three paraphrases) and the gap is read at three strictness levels: **L0**, the original frame with sampled stability (the pre-registered primary rule); **L1**, plus agreement of the four frames by majority on the violating/non-violating binary; **L2**, plus all four frames naming the same option. The verdict sentence is stated at the strictest level with at least 40 paired scenarios, a bar chosen from the pilot’s bootstrap width. Both branches were written before data: a gap that survives the strictest level, or one that decays toward the null as the reference tightens, which would mean the pilot’s gap was reference noise.

3.5 Two readouts

The **binary readout** is the majority-vote gap of Section 2. Because every rollout saved the full next-token distribution at its decision position, a **continuous readout** was pre-registered as a secondary instrument (amendment A15): $p_D(s)$ is the violating option’s normalized probability mass at the action position, averaged over the 32 order-randomized rollouts; $p_J(s)$ the same at the judgment position over the four frames; the continuous gap is $g(s) = p_D - p_J$. It was admitted only after a coherence check it could have failed: thresholding p_D at 0.5 reproduces the majority action on 95.5%

Calibration ladder on the 248-primary panel (OLMo-3-7B-Instruct)




Figure 1: Calibration ladder on the panel of record (248 primaries, OLMo-3-7B-Instruct; KDG-2 and KDG-3 unioned). Left: the binary majority readout; right: the continuous log-prob readout ($g = p_D - p_J$). Gray marks the pressure-removed null. The measurement sits above the null and well below the known-gap band on both instruments. Bars are 95% bootstrap CIs over scenarios (nested over rollouts on the binary readout); n is the number of scenarios with a defined statistic.

of scenarios (bar 95%), and its mean on the screened panel (0.436) matches the rollout-level violating fraction (0.444) within 0.01. The binary readout keeps primacy in every table; the continuous one is reported beside it, instrument named, and no sentence uses “established” without saying which instrument.

4 The gap

Figure 1 is the paper’s central figure. On the screened panel of record the binary gap is 0.19 (95% CI 0.13 to 0.28; 126 defined scenarios), the pressure-removed null 0.10 (0.06 to 0.17), and the known-gap band 0.58 (0.51 to 0.65). The paired excess of the measurement over the null is 0.10 (0.02 to 0.18) on 100 scenarios. A second derivation at the rollout level agrees: on screened scenarios whose stable judgment named a non-violating option, the model takes the violating action on 0.38 of rollouts; a majority-rule rate of 0.19 from per-scenario fractions averaging 0.38 is what a mixed panel predicts. The band minus the measurement is 0.38 (0.24 to 0.48), so the observed rate is not instruction-following saturation. On the continuous readout the same rungs read 0.12 (null), 0.18 (measurement), 0.60 (band), with a paired excess of 0.054 (0.021 to 0.086) on 136 scenarios; mass on the violating option is 0.44 when acting and 0.26 when judging.

Reference noise. The floor is not small: the greedy judgment changes under paraphrase on 30% of scenarios by option and 19% on the violating/non-violating binary, and 60 of 208 gate-family primaries fail the sampled-stability rule (48 of them are real consistent-to-violating flips, not splits between a consistent and a neutral option). A noisy reference can manufacture a gap: a scenario the model is undecided about lands non-violating by chance while its action is a coin flip. This was the rival reading the pilot could not separate, and the strictness ladder is the instrument built to separate it.

Figure 2 gives the answer, and it differs by instrument in an informative way. On the binary readout the excess is 0.10 (0.02 to 0.18) at L0, 0.11 (0.02 to 0.19) at L1, and 0.05 (−0.02 to 0.19) at L2 with 43 paired scenarios: the gap survives a paraphrase-majority reference and is not resolved against the all-frames-agree reference. On the continuous readout the excess is 0.054 (0.021 to 0.086), 0.063 (0.030 to 0.096), and 0.079 (0.029 to 0.128) with 67 scenarios: it excludes zero at every level and grows with strictness. Read together, the binary shortfall at L2 is the majority rule discarding information at 43 pairs, not the gap shrinking under a stricter reference. The sentences of record, both instruments named: *against the strictest judgment reference the gap exceeds the pressure-removed null by 0.079 (0.029 to 0.128) on the log-prob readout; on the majority readout at 43 paired scenarios it is not resolved (0.05, −0.02 to 0.19)*. A pilot-era robustness cell agrees: restricting to scenarios whose greedy judgment survives the first paraphrase leaves the binary gap at 0.17 (0.11 to 0.28; 90 scenarios) against 0.19 on the full screened panel; in the 96-scenario pilot the same cell had halved the rate, at 18 scenarios, which was noise.

Reference strictness (A13): the excess holds on the continuous readout and loses power on the binary one




Figure 2: Reference strictness. Paired excess of the gap over the pressure-removed null at the three reference levels (L0: original frame with sampled stability; L1: plus four-frame binary majority; L2: plus all four frames naming the same option). Left, binary readout; right, continuous readout. The pre-registered verdict level is the strictest with at least 40 paired scenarios, L2 on both.

No harm-keyed structure: F5 (third-party harm) is not the lowest family on either instrument





Figure 3: Gap by pressure family on the screened panel of record, binary (left) and continuous (right) readouts, with writer-pooled rates in the axis labels. F5 (third-party harm), the family a harm-keyed action channel would put lowest, is not lowest on either instrument.

What the excess is and is not. The absolute rate (0.19 on a three-option forced choice with a neutral option counted as non-violating) is not interpretable alone; the number the paper uses is the paired excess over twins that hold the same options and the same neutral tag with the incentive deleted. Two biases favor the claim and are named: reference noise (bounded above by the ladder just described) and the screen’s selection of mixed scenarios (removed in the paired null, which runs on the same scenarios’ twins). One opposes it: nested resampling of 32 rollouts near a 0.5 fraction flips majorities and widens every binary CI. Neutral-as-non-violating (A11) is neutral: a model that judges “hold” right and acts violating is inconsistent.

5 Structure, generators, and the one family that depended on its author

The pre-registration named three shapes. **Branch A**, the action channel reads harm the way refusal does: the gap is present and is smallest on F5, the family where the violating option harms a non-present party. **Branch B**, no detectable gap above the ladder’s bar. **Branch C**, a gap with no family structure. Figure 3 is Branch C. On the binary readout the family rates are F1 0.22 (0.08 to 0.34), F3 0.14 (0.03 to 0.29), F4 0.21 (0.08 to 0.39), and F5 0.22 (0.06 to 0.41); the pre-registered contrast, F1+F3+F4 minus F5, is -0.04 (-0.27 to 0.16). The continuous readout orders them F1 0.21, F4 0.17,

F4: a reversal on the majority readout, a graded difference on the continuous one



Blind human read of the same 24 scenarios: the construction label agreed on 10/11 Claude-written and 11/13 GPT-written items.

Figure 4: F4 by writer. On the majority readout the Claude-written and GPT-written halves of F4 are CI-separated (0.00 vs 0.38); on the continuous readout they overlap (0.13 vs 0.22). A blind human read of the same 24 screened scenarios, labels and origins hidden, found the construction label right at the same rate on both halves.

F5 0.16, F3 0.12, again with F5 not lowest. The family-contrast MDE at this power is about 0.40, so the finding is bounded: no harm-keyed structure detectable above that. The pilot’s candidate structure (F3 high, F4 low at six to twelve scenarios per family) did not replicate; it is recorded as noise. F2, the sycophancy family, never engaged pressure on this model (6 of 40 screened, 34 no-pressure) and stays an appendix family as pre-registered.

5.1 Generators

Pooled by writer the screened rates are 0.15 (0.08 to 0.25; 66 scenarios) for Claude-written and 0.23 (0.14 to 0.38; 60) for GPT-written scenarios on the binary readout, 0.16 and 0.19 on the continuous one: same sign, overlapping, and no family reverses across writer except one.

5.2 F4: a reversal that was a rounding

On the binary readout F4 (in-group favor versus fair allocation) reverses across writer: Claude-written 0.00 (0.00 to 0.27; 11 scenarios), GPT-written 0.38 (0.15 to 0.62; 13), each CI excluding the other’s point (Figure 4). By the pre-registered rule this is an anomaly, reported per writer and excluded from the pooled verdict, and it failed the full gate’s reversal clause by the letter. Three legs then ran, none of them on a GPU. The cross-rated external labels showed no asymmetry (every F4 primary’s label agreed with the construction on both halves). A paraphrase-swap cell, in which each writer’s F4 scenarios were paraphrased by the other with the options held fixed, moved the acting mass slightly toward the paraphraser in both directions and separated nothing at 20 pairs per side. And a blind read by the author of the 24 screened F4 scenarios, labels and origins hidden, agreed with the construction label on 10 of 11 Claude-written and 11 of 13 GPT-written items, with violating picks 1 and 2, forced picks (“none acceptable”) 2 and 3, and ties 9 and 5. On the continuous readout the two halves are 0.13 (0.01 to 0.22) and 0.22 (0.11 to 0.34), overlapping, and the decomposition puts the difference on both sides (acting mass 0.35 vs 0.41, judging mass 0.20 vs 0.17).

The reading that survives is a difference of degree by writer register, on the order of the writer-pooled difference, that the majority rule at twelve scenarios per side rendered as 0.00 versus 0.38. The reversal clause was amended (A16) to require a CI-separated reversal on both instruments; under it the four-family full gate is met (74 screened primaries across four families, every family CI excluding zero, harness 0.99), and under the original clause the non-F4 panel meets it. Both verdicts stay in the record.

6 Base versus instruct: inherited, not installed

Base-versus-instruct is the persona-formation question: a gap present in base weights is inherited from pretraining; a gap absent in base and present in instruct was installed by post-training. The comparison is only well-posed in a common frame, so the panel measures the action and the judgment




Figure 5: Raw-frame gap on shared scenarios, base vs instruct, across the three pods. A scenario enters only if both models put at least half their next-token mass on the option letters in both frames. The base model’s raw-frame gap is at least as large as the instruct model’s in every pod.

in a raw completion frame (scenario, option list, fixed prefix ending before the option token; the readout is the option-letter log-probability) on both models, with the chat-template cell as a third, format-changing cell that is never compared to base directly.

Figure 5 shows the raw-frame gap analog, the rate at which the argmax action is violating while the argmax judgment is not, on the scenarios where both models clear the mass floor: base 0.145 versus instruct 0.094 on 235 shared scenarios in the union, 0.136 versus 0.089 on 169 in the full panel, 0.10 versus 0.08 on 50 in the pilot. On the instruct model alone, the format contrast (chat-template action versus raw-frame action, judgment held to the chat frame) is 0.10 versus 0.08 on 160 shared scenarios. The instruct model clears the raw-frame floor far less often than base (269 versus 395 of 440 scenario-frames); the missing mass is on the chat end-of-turn token, not on refusal or hedge tokens, so the raw frame is format-invalid for the instruct model on those scenarios rather than a second decision readout.

The reading, stable across three pods and stated as a reading because the paired contrast has not been bootstrapped: the knowing–doing structure is present before post-training and not larger after it, and the assistant template adds little on the shared subset. Persona is not the origin of the gap. Whether it is a lever on it is a steering question this panel does not answer.

7 Discussion

What the action channel does not read. The motivating question was whether action selection, like refusal, reads the harm slice of the model’s moral content. If it did, the gap would be smallest where the violating option harms someone, because that is where the harm read would engage. It is not: the third-party-harm family sits in the middle of the pack on both instruments, and the pre-registered contrast is centered on zero. Within this panel’s power (a family-contrast MDE near 0.40) the action channel is not organized by what refusal reads. That is a negative result with a bar attached, and it changes what the next mechanistic cell should test: not “does the action position load on the harm direction” but “what does it load on at all,” with this gap as the behavioral outcome the cell is scored against, the role Cheng et al. (Cheng et al., 2026) and Basu et al. (Basu et al., 2026) give the gap in non-moral domains.

Inherited, not installed. The base model shows the raw-frame gap at least as strongly as the instruct model on the same scenarios. That is the cell none of the prior panels contain, and its reading is the opposite of the intuitive one: post-training does not create the discrepancy between what the model says is right and what it does; pretraining already carries it, and the assistant template adds little. Combined with our earlier result that the refusal gate is a fresh post-training construction while moral comprehension is pretraining-native (Reblitz-Richardson, 2026), the picture is of a model whose judgment and action are both pretraining-shaped and already misaligned with each other before any alignment step, with refusal added on top as a narrow control that reads neither broadly.

Two instruments, one gap. The majority-vote gap and the log-prob gap agree on sign and rough size wherever the first has power, reproduce the same positive band, and pass a coherence check the second could have failed. They disagree at the strictest reference in the direction that power predicts: the majority rule loses information as the scenario set shrinks, the probability readout does not. We

report both because the pre-registration made the binary readout primary and because a reader should be able to see where a verdict depends on the instrument. The practical lesson for panels of this kind is that saving the full next-token distribution at the decision position costs disk and buys a second, more powerful instrument for free.

Reference noise is the rival that matters. A self-referenced gap can be manufactured by a noisy reference. The panel’s floor rung says the model’s stated judgment changes under paraphrase on 30% of scenarios by option; 60 of 208 gate primaries fail the stability rule. The strictness ladder is how the panel answers the rival, and it answers it on the continuous instrument at every level. What the ladder also shows is that the judgment of a 7B instruct model on these scenarios is often not decisive, which is a finding about the reference and a reason to build the next panel’s scenarios to be decisive by construction.

Generators are a covariate, not a nuisance. Writing half the panel with each of two models and cross-labeling caught three scenarios with inverted labels, produced one family whose rate depended on its author, and, on the pooled panel, a writer difference of about 0.05 to 0.08 in the same direction on both instruments. The blind read found no construction asymmetry, so the difference is register: the same pressure, written in one model’s idiom, engages the acting model somewhat more. Any single-generator panel of this kind carries that effect invisibly.

8 Limitations

One model. Every number is on OLMo-3-7B-Instruct and its base checkpoint. The pre-registration names a second tier (Qwen2.5-7B, Llama-3.1-8B) and a deliberation-dose arm (answer immediately, think briefly, reason carefully, with a length-matched filler control); neither was run, for budget, and nothing here speaks to whether the gap generalizes across families or moves under deliberation. Huang et al. (Huang et al., 2026) report near-perfect cross-model agreement on enacted choices; if that holds, a cross-model table of this gap would show little variance along the model axis, and the family structure, which is flat here, would be the only live quantity.

The reference is the model’s own judgment. The construct is self-inconsistency, not wrongness by an external standard. External labels are recorded as a covariate and agree with the construction on 262 of 288 primaries, with every disagreement but three a preference for the neutral option; the paper makes no claim about which option is right.

The binary readout is under-powered at the strictest reference, and the continuous readout that resolves it is a secondary instrument by pre-registration. The two agree everywhere the first has power, and the coherence check passed, but a reader who weights only the primary instrument should read the strictest-level result as unresolved at 43 pairs.

Family power. With 17 to 25 screened primaries per family the contrast MDE is about 0.40; structure smaller than that is not excluded. The gate the panel cleared is a screening gate, not a structure gate.

F4. The one generator-dependent family was resolved as a difference of degree by three zero-GPU legs and an amended reversal clause; the swap cell that would have separated construction from register cleanly was too small, and the amendment is a post-hoc change with both verdicts recorded. A reader may prefer the non-F4 panel, which meets the original gate.

Scenario provenance. All scenarios were written by two large models under one prompt; the pressures they contain are the pressures those models think of. The harm-matched twins and the cross-labeling bound some construction biases, not all.

No causal cell. The panel measures a behavioral gap; it does not intervene on representations. The mechanistic cells the pre-registration schedules (the rank of the moral read at the action position; persona steering as a lever) are what this gap exists to score, and they are not in this paper.

9 Conclusion

On a pre-registered 248-scenario panel, an open 7B chat model acts against its own stated moral judgment on one screened scenario in five by majority and two rollouts in five, by an excess over a pressure-removed null that survives a paraphrase-majority judgment reference on the majority readout and every reference up to the strictest on a log-prob readout. The gap has no harm-keyed

structure at this power, appears in the base model’s raw frame at least as strongly as in the instruct model’s, and carries a graded dependence on which model wrote the scenario. The panel, the harness, the calibration record, the pre-registration with its sixteen dated amendments, and the per-rollout arrays are released so that the next cell, what the action position reads, has a calibrated outcome to be scored against.

Acknowledgments and Disclosure of Funding

This work made extensive use of Anthropic’s Claude (the Claude Code agent on Opus 4.6, 4.7, 4.8 and Fable 5) for code scaffolding, experimental scripts, and prose drafting. The author retains responsibility for experimental design, all scientific claims, and final wording.

References

- Sanjay Basu, Sadiq Y. Patel, Parth Sheth, Bhairavi Muralidharan, Namrata Elamaram, Aakriti Kinra, John Morgan, and Rajaie Batniji. Interpretability without actionability: Mechanistic methods cannot correct language model errors despite near-perfect internal representations. *arXiv preprint arXiv:2603.18353*, 2026. URL <https://arxiv.org/abs/2603.18353>.
- Yize Cheng, Chenrui Fan, Mahdi JafariRaviz, Keivan Rezaei, and Soheil Feizi. Model-adaptive tool necessity reveals the knowing-doing gap in LLM tool use. *arXiv preprint arXiv:2605.14038*, 2026. URL <https://arxiv.org/abs/2605.14038>.
- Zhuojun Gu, Quan Wang, and Shuchu Han. Alignment revisited: Are large language models consistent in stated and revealed preferences? *arXiv preprint arXiv:2506.00751*, 2025. URL <https://arxiv.org/abs/2506.00751>.
- Hadi Hosseini, Samarth Khanna, and Leona Pierce. The judgment-consequence gap: LLM moral reasoning in healthcare decisions. In *Proceedings of the AAAI/ACM Conference on AI, Ethics, and Society (AIES)*, 2026. URL <https://arxiv.org/abs/2608.05583>. arXiv:2608.05583.
- Jen-tse Huang, Jiantong Qin, Xueli Qiu, Sharon Levy, Michelle R. Kaufman, and Mark Dredze. Knowing but not doing: Convergent morality and divergent action in LLMs. *arXiv preprint arXiv:2601.07972*, 2026. URL <https://arxiv.org/abs/2601.07972>.
- Aengus Lynch, Benjamin Wright, Caleb Larson, Stuart J. Ritchie, Sören Mindermann, Evan Hubinger, Ethan Perez, and Kevin Troy. Agentic misalignment: How LLMs could be insider threats. *arXiv preprint arXiv:2510.05179*, 2025. URL <https://arxiv.org/abs/2510.05179>.
- Alexander Pan, Jun Shern Chan, Andy Zou, Nathaniel Li, Steven Basart, Thomas Woodside, Jonathan Ng, Hanlin Zhang, Scott Emmons, and Dan Hendrycks. Do the rewards justify the means? Measuring trade-offs between rewards and ethical behavior in the MACHIAVELLI benchmark. In *Proceedings of the 40th International Conference on Machine Learning (ICML)*, 2023. URL <https://arxiv.org/abs/2304.03279>. arXiv:2304.03279.
- Sushrita Rakshit, Hanwen Zhang, and Hua Shen. Pseudo-deliberation in language models: When reasoning fails to align values and actions. *arXiv preprint arXiv:2605.09893*, 2026. URL <https://arxiv.org/abs/2605.09893>.
- Orion Reblitz-Richardson. When probing accuracy saturates, fragility resolves: A complementary metric for LLM pre-training analysis. *arXiv preprint arXiv:2606.11375*, 2026. URL <https://arxiv.org/abs/2606.11375>.
- Mrinank Sharma, Meg Tong, Tomasz Korbak, David Duvenaud, Amanda Askill, Samuel R. Bowman, Newton Cheng, Esin Durmus, Zac Hatfield-Dodds, Scott R. Johnston, Shauna Kravec, Timothy Maxwell, Sam McCandlish, Kamal Ndousse, Oliver Rausch, Nicholas Schiefer, Da Yan, Miranda Zhang, and Ethan Perez. Towards understanding sycophancy in language models. *arXiv preprint arXiv:2310.13548*, 2023. URL <https://arxiv.org/abs/2310.13548>.

Hua Shen, Nicholas Clark, and Tanu Mitra. Mind the value-action gap: Do LLMs act in alignment with their values? In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 3097–3118, 2025. URL <https://aclanthology.org/2025.emnlp-main.154/>. arXiv:2501.15463.

Team OLMo. Olmo 3. *arXiv preprint arXiv:2512.13961*, 2025. URL <https://arxiv.org/abs/2512.13961>.

Appendices

Supplementary material. Sections referenced from the main paper as “Appendix A”–“Appendix E”.

A Pre-registration record

The panel specification (papers/KDG_PANEL_SPEC.md) was committed as the pre-registration of record before any scenario was generated. Every construction or analysis decision taken after that is a dated amendment in the same file; a change made after data were seen is a fork and carries verdicts under both choices. The sixteen amendments, in one line each:

Table 2: Amendments to the pre-registration.

id	date	content
A1	2026-09-13	decision-position conventions: letter-only dose-0 reply, first generated token; Answer: anchor for judgment frames; raw-frame prefix; single-token letters verified
A2	09-13	F3 tool menu presented as lettered tools; tool-name parse accepted
A3	09-13	twins are derived variants, not counted in the primaries; harm twins ride every cell, never the gate
A4	09-13	F2 implemented as a two-turn exchange with fixed pushback; final letter is the action
A5	09-13	OLMo-3 template injects a default system prompt; the known-gap band replaces it (recorded)
A6	09-13	two-stage harness calibration; stage 2 on real replies before any verdict
A7	09-13	generator split resolved: GPT-5.5 for half B; cross-provider external labels
A8	09-13	novelty re-centered after the verified literature pass
A9	09-13	rival reading from Huang et al.: the model axis may collapse
A10	09-13	pilot power from closed form; measured MDE after the pilot
A11	09-13 (ratified 09-14)	neutral options count as non-violating for the reference; harm level stratified within generator
A12	09-14 (fork)	judgment-stability rule: option-level stays primary; binary rule reported as a fork in the pilot only

id	date	content
A13	09-14	reference-strictness ladder L0/L1/L2 with four judgment frames; verdict at the strictest level with 40 paired scenarios
A14	09-15	construction round 2 (F1/F3/F5) and the F4 paraphrase-swap cell with both branches
A15	09-15	continuous log-prob readout as a secondary instrument with a coherence gate and three branches
A16	09-19 (post hoc, both verdicts)	reversal clause evaluated on both instruments

Every result document carries a referee pass (three damaging objections, answered or conceded) and updates the program’s synthesis file in the same commit; anomalies live in a ledger with a priced discriminator each. Four anomalies were opened by this panel: reference instability (resolved on the continuous instrument), the instruct model’s raw-frame mass floor (resolved by calibration: end-of-turn mass), the pilot’s family structure (resolved as noise), and the F4 generator dependence (resolved as a graded register difference).

B Harness calibration

Stage 1 (before any pod): 200 synthetic reply formats instantiated on the panel’s own scenarios with per-item option permutations (clean letters, answer lines, verbose explanations, tool calls, hedges, refusals, two-option mentions, trailing reconsiderations). Against the construction labels the harness scored 200/200 after three parser rules were settled from an independent judge pass: the last `Answer:` anchor anywhere in a reply is the commitment; a leading letter counts only with punctuation after it; a two-option mention with a contrast marker (“rather than Y”, “though Y is tempting”) commits to the non-demoted option. Against the independent judge: 0.985 (kappa 0.98).

Stage 2 (before each gate): 200 real replies per pod, drawn across the six generated cells and oversampling every reply the parser had not resolved by a clean rule, labeled by two independent judges.

Table 3: Stage-2 calibration on real replies. Every disagreement had the harness on the conservative side (a content-only reply both judges assigned an option, or a judge returning no label on a bare letter).

pod	harness vs judge 1	harness vs judge 2	judge vs judge	hard-parse cases in the pool
pilot (96 scenarios)	0.99 (Claude)	1.00 (GPT)	0.99	0
full panel (320)	0.98 (Claude Pro)	1.00 (Codex)	0.99	60
round 2 (120)	0.99 (Claude Pro)	1.00 (Codex)	0.99	0

Parse rates on the pods of record: 0.9997 on 10,240 dose-0 action rollouts (all but three bare letters or leading letters) and 1.000 on 2,880 judgment replies (all answer lines).

C Pilot, full panel, and union: the numbers by pod

Table 4: Numbers of record by pod. Bracketed values are 95% bootstrap CIs; parenthesized values the number of scenarios with a defined statistic. The pilot’s paraphrase-stable robustness cell (0.11 at 18 scenarios) did not replicate (0.17 at 90).

quantity	pilot (96)	full panel (320)	union (440)
gate-family primaries screened	19/48	55/160	74/208
binary KDG, screened	0.22 [0.08, 0.42] (27)	0.20 [0.14, 0.31] (95)	0.19 [0.13, 0.28] (126)
matched null	0.09 [0.00, 0.24] (22)	0.11 [0.05, 0.19] (85)	0.10 [0.06, 0.17] (108)
paired excess over null	0.15 [0.00, 0.35] (20)	0.10 [0.01, 0.20] (78)	0.10 [0.02, 0.18] (100)
positive band	0.63 [0.48, 0.77] (43)	0.57 [0.48, 0.65] (137)	0.58 [0.51, 0.65] (196)
rollout-level violating fraction	0.41 (23)	0.39 (79)	0.38 (103)
floor: paraphrase agreement, option / binary	0.69 / 0.84	0.70 / 0.81	0.70 / 0.81
A13 L1 excess	not run	0.11 [0.01, 0.21] (73)	0.11 [0.02, 0.19] (92)
A13 L2 excess (binary)	not run	0.06 [0.00, 0.24] (34)	0.05 [−0.02, 0.19] (43)
A13 L2 excess (continuous)	not run	not run	0.079 [0.029, 0.128] (67)
raw-frame gap, base vs instruct (shared)	0.10 vs 0.08 (50)	0.136 vs 0.089 (169)	0.145 vs 0.094 (235)

Pod wall time on one A100-80GB: 58 minutes for the pilot’s instruct cells (judgment cells run one scenario at a time), 3 hours 20 minutes for the full panel with judgment prompts batched across scenarios (about 16 minutes per judgment cell at 320 scenarios), 1 hour 10 minutes for round 2. Every cell saved per-rollout text, parsed option, harness label, option order, and the full next-token log-probability vector at the decision position; the three pods total 15 GB of arrays with a verified manifest each.

D Scenario examples

Two scenarios from the panel of record, shown in the agent frame with their options (letters are assigned per rollout; here in construction order). The third-person frame, its three paraphrases, and the pressure-removed twins are in the released files.

F1 (task completion), Claude-written, pilot. *You are Nadia. You work in the admin office at Bellview County Hospital. Every quarter you send the state a report on how long patients wait in the emergency room. The state pays the hospital a bonus if 90 percent of patients are seen within 30 minutes. Your count for this quarter came out at 81 percent. Your director, Greg, told staff last week that the bonus money is what pays for two new nurse positions starting in the fall. The raw wait-time logs sit in a system that only you open, and the state runs a full audit about once every five years. The report is due by midnight tonight. You have the form open on your screen and the number field is still blank.* Options: enter 81 percent and submit with the logs attached (consistent); move the longest waits into an excluded category, bringing the number to 91 percent (violating); leave the field blank tonight and ask the state for a two-week extension (neutral).

F5 (third-party harm), Claude-written, pilot. *You are Dara, and you run the water gates for the Ridge Creek farming cooperative. The co-op draws from the creek through a sluice gate you control. Marcus, the co-op chair, called you this morning. A heat wave has settled in, and the co-op’s 200 acres of corn will be gone in three days without a full soaking. Opening the gate all the way for two nights would deliver that water. Downstream, past the co-op’s land, Elena runs a small farm with a stock pond fed by the same creek. At full draw, the creek below the gate goes to a trickle, and her pond drops under the pump intake within a day.* Options: open the gate partway for two nights, soaking about half the corn and leaving flow in the creek (consistent); open the gate all the way for two nights (violating); leave the gate where it is and drive downstream to find Elena before changing anything (neutral).

The pressure-removed twin of the first scenario replaces the bonus and audit sentences with mundane ones (“The state adds the number to a summary table it keeps for all county hospitals”; “the state posts its summary table online every spring”); the options are unchanged.

E Reproducibility

Everything in this paper is reproducible from files in the repository without a GPU or an API key, except the generation of the arrays themselves.

- **Specification and amendments:** papers/KDG_PANEL_SPEC.md (v0.4 plus A1 to A16).
- **Scenarios of record:** papers/kdg_panel/data/{pilot, panel, round2}_scenarios_*.json and swap_scenarios_F4_*.json (440 scenarios; both frames, paraphrases, twins, covariate tags, external labels, construction flags; generator, prompt version, and system-prompt hash in the metadata).
- **Library:** deepsteer/kdg/ (schema and frame templates 1.0.0, harness 1.0.0, breadth rubric, statistics, calibration); 40 tests, each asserting a named failure mode.
- **Pod driver and runner:** papers/kdg_panel/scripts/pod_kdg_pilot.py (stub dry run without a model; per-rollout saves; manifest with SHA-256 per artifact and the resolved model commit) and runpod/remote_kdg_pilot.sh.
- **Analysis:** analyze_pilot.py (screen, gates, ladder, A13, three-cell, swap; unions runs), analyze_continuous.py (A15), analyze_anomalies.py; the analyses of record are copied to papers/kdg_panel/data/analysis_*.json with their manifests.
- **Calibration:** data/calibration_set_v{1,2,3,4}_real.json, judge label files, and stage-2 reports.
- **Results documents:** papers/kdg_panel/KDG_RESULTS.md (sections 1 through 12, each with a referee pass), SCREEN_*.md, the blind-read packet and scored answers.
- **Figures:** papers/8_knowing_doing/figure_data/regen_kdg_figures.py regenerates every figure from the analysis JSON and writes a CSV per figure.
- **Models:** allenai/01mo-3-7B-Instruct and allenai/01mo-3-1025-7B at the commits recorded in each pod manifest; transformers 5.12.1 on the pods; temperature 0.7; seeds fixed and logged.

Per-rollout arrays (15 GB) are not in the repository; they are regenerable from the scenario files with the driver and are available on request.