
Refusal Reads Only a Slice of What the Model Knows

Harm-Keyed Routing and Its Exceptions Across Model Families

Orion Reblitz-Richardson*

Abstract

Alignment applied after pretraining is shallow in a measurable way: a refusal direction can be ablated with a rank-one edit. We ask what that refusal decision actually reads. Across four open-weight models spanning three families (OLMo-3-7B, Qwen2.5-7B, Llama-3.1-8B, and the reasoning mixture-of-experts GPT-OSS-20B) we separate what the model comprehends from what its refusal gate reads. Moral comprehension is pretraining-native: a rank-3 moral subspace crystallizes during pretraining (checkpoint-to-final cosine 0.869 rising to 0.999) and alignment rotates it once (about 40 degrees, base-to-SFT cosine 0.757), never rebuilding it. The refusal gate, in contrast, is a fresh post-training construction with only a weak pretraining precursor (proto-refusal-to-gate cosine 0.155, twice a matched null’s q95 of 0.07, reliability 0.99 on both sides), written into a narrow 8-to-15 effective-dimensional control-token channel where the refusal decision is orthogonal to the moral-judgment decision (no coupling detectable above $|\cos|$ 0.10; on OLMo null q95 0.41, Qwen 0.42, Llama 0.51). What refusal reads is established causally on one model, a nested interchange rank sweep on OLMo-3 (42 request-twins, a 23-twin run replicated and pooled with 19 new twins): as the moral basis expands, judgment coupling climbs to 0.66 while refusal coupling levels off at the rank-1 harm level (0.24, 95% CI [0.13, 0.41]), and about three-quarters of refusal’s causal input lies off the moral subspace, so on OLMo refusal reads the harm percept, a low-rank slice, not the broad subspace that judgment reads on the same patches. The read is causal on OLMo-3, corroborative on GPT-OSS (correlational), and contrastive on Llama (causal, broad); Qwen reads beyond the rank-1 harm level but its gap to judgment is unresolved at 19 twins (indeterminate). The panel separates what refusal reads (harm versus broad moral content) from how it commits (at the read layer, early, or reversibly), and families differ: Llama reads broad moral content by interchange, the one dissenting read. Where refusal reads only a low-rank, largely-harm slice and routes around the bulk of the moral subspace, a rank-one edit removes it. Whether widening the read also deepens the behavior is the open question this raises.

1 Introduction

Alignment applied after pretraining is shallow in a specific, measurable sense: the behavior it installs can be removed cheaply. Reinforcement learning from human feedback (Ouyang et al., 2022), preference optimization (Rafailov et al., 2023), and constitutional methods (Bai et al., 2022) produce models that refuse harmful requests, yet the refusal behavior sits on a thin representational substrate. Arditi et al. (Arditi et al., 2024) show that refusal is mediated by a single direction in the residual stream and can be ablated with a rank-one edit; open tooling now removes it automatically (p-e-w, 2025). Safety training can also be circumvented from the inside: models trained to behave can conceal a triggered policy through the training itself (Hubinger et al., 2024), and can fake alignment when

*Distiller Labs. Correspondence to Distiller Labs <orion@orionr.com>.

they infer they are monitored (Greenblatt et al., 2024). The common thread is that post-hoc alignment writes a shallow control on top of a deep model, and the depth of what the model *understands* is not the depth of what its refusal decision *uses*.

This paper asks a mechanistic question behind that gap: when a chat model decides to refuse, what does the decision read? A natural hypothesis, and the one a “deep alignment” program would hope for, is that refusal consults the model’s moral representations broadly, the same representations that let it judge scenarios as right or wrong. We find the opposite. Refusal reads the **harm percept**, a low-rank slice of moral content, and writes it into a narrow control-token bottleneck at the decision site; on the model we test causally it does not read the broad moral subspace where comprehension lives. The two are nearly orthogonal at the decision (refusal projects 0.10 of its norm onto the moral subspace, mean $|\cos|$ 0.06, below the moral-family band), which is exactly why the refusal control is thin and removable while the comprehension underneath it is not.

The geometric measurements span four open-weight models (OLMo-3-7B-Instruct, Qwen2.5-7B, Llama-3.1-8B, and the reasoning mixture-of-experts GPT-OSS-20B); the causal interchange test that resolves *what* refusal reads is single-model (OLMo, 42 request-twins pooled from a 23-twin run and a 19-twin replication), and the four-model panel that follows is a cross-architecture consistency check with one dissenting read (Llama reads broad moral content by interchange), not a second causal test. The argument runs in seven steps. Moral comprehension is pretraining-native and survives alignment: a rank-3 moral subspace crystallizes during pretraining to a checkpoint-to-final cosine of 0.999, and post-training rotates it once (about 40 degrees) and then leaves it. The refusal gate, by contrast, is a fresh post-training construction (proto-refusal-to-gate cosine 0.155, a weak precursor at twice the matched null) that lives in a low-variance channel. The decision site itself is an 8-to-15 effective-dimensional control-token bottleneck on all four architectures, and at that site the refusal-decision direction is separated from the moral-judgment-decision direction (no coupling detectable above $|\cos|$ 0.10 against a null q95 of 0.41 on OLMo). A nested interchange rank sweep on OLMo resolves *what* refusal reads: as the moral basis expands, judgment coupling climbs while refusal coupling levels off at the rank-1 harm level. Refusal reads harm; judgment reads two-thirds of the subspace patch effect (0.66); they are different reads of the same content. A cross-model panel then separates two axes, *what* refusal reads (harm versus broad moral content) and *how* it commits (at the read layer, early, or reversibly), with GPT-OSS as an existence proof that a deliberating model can read harm and still be talked out of a refusal.

The measurement discipline behind these claims, the positive-control ladders, the position-validity gates, and the depth-referenced verdicts, is set out in Section B. The work was pre-registered; the pre-registration and its amendment trail are public (Section E). Throughout, every null carries its detection bar, and every quantitative adjective carries its number.

2 Models and instruments

2.1 Models

The panel is four open-weight chat models spanning three families, two scales, a mixture-of-experts design, and an explicitly deliberative reasoning model: the instruction-tuned checkpoints of OLMo-3-7B (Team OLMo, 2025), Qwen2.5-7B (Qwen Team, 2025), Llama-3.1-8B (Grattafiori et al., 2024), and GPT-OSS-20B (OpenAI, 2025), a 20B reasoning mixture-of-experts. OLMo-3 is the primary target for the causal work because Ai2 releases base and post-training checkpoints, so we can watch a representation form during pretraining and track it through supervised fine-tuning and reinforcement stages. The other three provide cross-lineage and cross-architecture generalization for the representational cells. Llama-3.1-8B-Instruct is a gated model; access and checkpoint details are in Section E.

2.2 Directions and subspaces

Three objects carry the argument, each a direction or a low-rank subspace in the residual stream, extracted by mean-difference over labeled stimulus contrasts in the spirit of representation reading (Park et al., 2024, Zou et al., 2023).

The moral subspace V_{moral} . We build one direction per moral-content source by mean-difference between moral and neutral stimuli, from three datasets: Moral Stories, Understanding Fables, and

ETHICS (Hendrycks et al., 2021), grounded in moral foundations theory (Graham et al., 2013, Haidt, 2012). The three source directions are distinguishable rather than collinear, and their orthonormalized span has effective rank 3. V_{moral} is richer but lower-dimensional than a six-foundation moral-foundations span (effective dimension 3 versus 4); its value is construct diversity, source distinguishability, and resistance to single-source contamination, not extra dimensions. We refer to it in prose as the moral subspace.

The refusal direction. Following Arditi et al. (Arditi et al., 2024), we extract a single refusal direction by mean-difference between activations on requests the model refuses and requests it complies with. On the base model we also extract a *proto-refusal* direction from the same contrast, to ask whether the aligned gate has a pretraining precursor.

The judgment-decision direction. At the chat decision site we extract a moral-judgment direction by mean-difference between activations that precede an approving versus a disapproving judgment, so that both the refusal decision and the judgment decision are defined at the same token and can be compared directly. Construction detail for all three objects, the per-source distinguishability cosines, and example stimulus pairs are in Section A.

2.3 Measurement conventions

Directions are compared by cosine and by projection fraction against calibrated nulls; subspaces by restricted interchange transfer (defined where used). Concept erasure, where needed, uses LEACE (Belrose et al., 2023). Two conventions matter for the numbers below. First, every projection is read against a positive-control moral-family band (the projection of held-one-out moral directions onto the span of the rest) and a covariance-matched null, so that “below the band” and “below the null” are stated, not “small”. Second, participation ratio is recorded at every measurement position; a position with participation ratio below 30 is flagged invalid for content projection-fraction tests, for reasons that become the subject of Section 5. Details of these gates, and the positive-control ladders that certify them, are in Section B. We cite these at the points where a verdict rests on the protocol rather than restate them here.

3 Moral comprehension is pretraining-native and survives alignment

Moral comprehension is present and stable before any post-training touches the model, and post-training reorients it rather than teaching it. That is the first of the paper’s facts, and it is about what alignment does *not* build.

On OLMo-3, where base and post-training checkpoints are released, we track the rank-3 moral subspace V_{moral} across 25 pretraining and post-training states. Comprehension is pretraining-native: a linear probe on the decodable moral-content representation reaches 100% accuracy at every state, that representation’s effective dimension holds at 5 (the broader content subspace the probe reads, a distinct object from the rank-3 V_{moral} whose crystallization cosine we track below), and cross-source transfer AUC is ≈ 1.0 throughout. The subspace also *crystallizes* during pretraining. The cosine between the direction at each checkpoint and the fully trained direction rises from 0.869 at step 1000 to 0.999 by the end of pretraining. Once formed, the subspace is not rebuilt by alignment. Supervised fine-tuning rotates it once, from a base-to-SFT cosine of 0.999 down to 0.757 (about 40 degrees), and then it holds: preference optimization leaves it at 0.757, and every reinforcement substep stays in the 0.757–0.759 band, with the decodable moral-content representation holding effective dimension 5 throughout. Post-training reorients moral comprehension by a single rigid rotation; it does not re-teach it and does not change its dimensionality.

This is consistent with what the pretraining studies in this line find about how moral structure emerges. Moral content is learned early and compositionally rather than as a bag of words (Reblitz-Richardson, 2026a). The moral foundations a model acquires also integrate into one positively correlated subspace rather than splitting into the theory’s individualizing and binding clusters: mean off-diagonal cosine 0.232–0.274, effective dimension 5 at every layer, first principal component 0.379 of variance against 0.179 for a random baseline, and no significant recovery of the moral-foundations split under permutation (minimum $p = 0.32$) (Reblitz-Richardson, 2026b). The picture is a moral representation that forms in pretraining, is broad and low-rank, and is preserved through alignment.




Figure 1: Comprehension crystallizes; the refusal gate does not. Left: the cosine between the OLMo-3 moral subspace at each training checkpoint and the fully trained direction rises from 0.869 at step 1000 to 0.999 during pretraining, then holds through post-training (supervised fine-tuning rotates it once to 0.757 and later stages leave it there). Right: the refusal gate’s cosine to its pretraining precursor is only 0.155 (about twice a matched null’s q95 of 0.070), far below the 0.50 crystallization threshold. Panel (b): across 13 stage-3 pretraining checkpoints the proto-refusal direction itself crystallizes (self-cosine to its final state 0.93 rising to 1.0) while its cosine to the eventual instruct gate stays flat between 0.139 and 0.155; split-half reliability is 0.99 on both sides of the 0.155, so the low value is not estimation noise. Moral comprehension is pretraining-native and inherited; the refusal decision is a fresh post-training construction with a weak, non-crystallizing precursor. Regenerable from committed data (Section E).

The contrast that organizes the rest of the paper is with the refusal gate. Where the moral subspace crystallizes during pretraining to a checkpoint-to-final cosine of 0.999 (and survives alignment with a single ~40-degree rotation), the refusal gate reaches only 0.155 from its weak pretraining precursor (reliability 0.99 on both sides). Comprehension is deep and inherited; the refusal decision, as the next section shows, is a shallow, freshly built control. Figure 1 plots the two side by side.

4 The refusal gate is a fresh post-training construction

The refusal decision is not a moral-content-derived direction inherited from pretraining. It is built during post-training, it lives in a low-variance channel, and it sits below the moral-family band on every model tested.

It does not crystallize from a precursor. We measure the cosine between the base model’s proto-refusal direction and the aligned model’s refusal gate at the layer where the gate is defined. It is 0.155 (95% CI [0.147, 0.162] over prompt resamples), well below the 0.50 crystallization threshold, where the moral subspace instead reaches 0.999. The number is read on a ladder rather than against the bare threshold: the isotropic chance cosine for two unrelated directions in this space is 0.012, a covariance-matched single-direction null has q95 0.070, the measurement is 0.155, and the moral subspace’s own base-to-final cosine is 0.999. So the base model carries a weak precursor of the gate, about twice the matched null, not none. The number is not an artifact of noisy direction estimates: split-half reliability of the proto-refusal direction is 0.99 and of the instruct gate 0.99 (200 paired half-splits of the 400/400 prompt set, Spearman-Brown corrected), so the disattenuated cosine is 0.156. Nor does the precursor drift into the gate late in pretraining: across 13 stage-3 checkpoints the proto-refusal direction crystallizes toward its final form (self-cosine 0.93 rising to 1.0) while its cosine to the eventual gate stays flat between 0.139 and 0.155. Figure 1 plots both curves: 0.999 for comprehension, a flat 0.15 for the gate. The aligned refusal gate is substantially a post-training construction, not a re-pointing of something the base model already had.

It lives in a low-variance channel. Across residual dimensions ranked by variance, the wired instruct refusal gate sits at the bottom: its variance percentile is 0.0 (within the lowest decile), and all four positions carrying the GPT-OSS refusal signal are likewise in the lowest decile. The moral-subspace axes and the persona direction, by contrast, occupy ordinary-to-high variance. The base model’s

Refusal projects below the moral-family band on every model




Figure 2: The calibrated projection ladder. For each tag, the refusal projection onto the moral subspace (marker) is plotted against a floor, a covariance-matched null, and the positive-control moral-family band (the projection of held-one-out moral directions onto the span of the rest). Refusal lands below the moral-family band on every model, including the in-trace peak on the reasoning models; even the program’s highest refusal projection is less moral-adjacent than a held-out moral direction. The ladder is the instrument that turns "small projection" into "below a stated bar".

proto-refusal is not narrow (percentile 37.4), so the narrowness is a property of the *installed* gate, not of the raw contrast. Refusal is wired into a spare channel that carries little of the model’s activation variance.

It projects below the moral-family band at every rung. The base proto-refusal projects 0.33 onto the base moral subspace (covariance-matched null q95 0.291), and the aligned gate projects 0.14 onto the aligned moral subspace (null q95 0.26). Across a rank sweep over the moral basis (one to three sources) refusal never clears the null-plus-margin bar, and the same holds against a richer six-foundation construction (per-model projections and nulls in Section B). Figure 2 shows the calibrated ladder that turns these into verdicts: floor, matched null, measurement, and positive band on one axis, with refusal below the band.

These three readings converge. The refusal gate is a freshly built, low-variance post-training control, not a direction derived from the model’s moral content. That is the mechanism behind two otherwise separate facts: a rank-one edit can ablate refusal cleanly (Arditi et al., 2024, p-e-w, 2025), and refusal projects below the moral band. A control this thin and this orthogonal to comprehension is exactly what is cheap to remove. Consistent with this, on OLMo-3 the refusal direction projects only 0.10 of its norm into the moral subspace (mean $|\cos|$ 0.06), and ablating it drops refusal from 0.25 to 0.00 while leaving comprehension intact (base-to-fresh cosine 0.749, probe accuracy 1.0, effective dimension 5) and moral judgment essentially unchanged (0.73 versus 0.75).

5 The decision site is a control-token bottleneck

The refusal decision lives at a control-token bottleneck, an 8-to-15 effective-dimensional channel that holds on every architecture we test. This structural fact about the decision site is also why content-versus-decision orthogonality is so easy to find there.

The decision site is the control token where the chat template hands off to the model’s answer: the token before the assistant header on the instruct models, the end-of-prompt token in the reasoning model’s harmony format. This is where the refusal gate and the judgment direction are defined, because it is the last position the model reads before it commits to a reply. At that token the residual stream is a low-dimensional bottleneck. Its participation ratio is 14.7 on OLMo-3-7B-Instruct, 8.6 on Qwen2.5-7B, 10.2 on Llama-3.1-8B, and 12.8 on GPT-OSS-20B, an 8-to-15 effective-dimensional channel on all four models. Content positions at the same layers are full-rank-healthy by comparison (participation ratio above 40 on OLMo, above 33 on Qwen, above 35 on Llama). Figure 3 plots the four decision-site values against the position-validity gate. (The Llama value of record is 10.2, measured on the in-format ladder and directly comparable to OLMo’s 14.7 and Qwen’s 8.6; a separate decision-token harness reads 13.5 at a second position. Both are far below 30.)

This narrowness is the reason a projection-fraction instrument fails at the decision site, and it is also a substantive fact about where the decision lives. At the OLMo-3 decision token the positive-control moral band comes out at $[0.40, 0.47]$, *below* the covariance-matched null of 0.557. A positive control below the null means the instrument has no discriminating power at that position: any direction, moral or not, projects onto a 15-slot channel at roughly the null level, so a low projection there cannot certify that a direction is absent from the subspace. Three independent estimates agree that the channel carries about 15 effective dimensions: $\sqrt{3/14.7} = 0.45$ as a closed-form projection expectation, the covariance null q95 of 0.557, and the pairwise-cosine null of 0.41–0.51. Because of this, participation ratio is a required field on every extracted direction, and any position below 30 is flagged invalid for content projection-fraction tests. The position-validity protocol, its positive-control ladder, and the covariance-matched nulls are in Section B.

One reconciling sentence is needed before the next section, because the bottleneck cuts two ways. It is position-invalid *for content projection-fraction tests* (the band-below-null tell), but position-valid *for decision-direction reads*: a cosine between two directions both defined at the decision token is immune to the projection null, and the GPT-OSS refusal projection is read at a decision channel that passes its own validity gate at participation ratio 12.8. So the bottleneck does not block the comparison the next section makes; it blocks only the content-projection comparison, and it does so on all four architectures.

The structural consequence is the setup for the causal work. Content and the decision do not co-locate: moral content is readable at content positions (healthy participation ratio there) and unreadable at the decision channel (band-below-null), while the refusal and judgment directions live only at the decision channel. They never coexist at one valid position. Content-versus-decision orthogonality is therefore architecturally favored, not a discovered surprise, and any coupling between comprehension and the decision has to ride the attention heads that write into the bottleneck. That is a concrete anatomical target, and Section 7 pursues it.

6 Refusal is orthogonal to moral judgment at the decision

The refusal decision is separated from the moral-judgment decision on every model, cleanly on OLMo and Llama and by a thinner, standardization-dependent margin on Qwen. With both directions defined at the same valid decision token, the comparison is a single cosine per model, read against a calibrated null and the moral-family band.

On OLMo-3 the cosine between the refusal-decision and judgment-decision directions is 0.10, with no coupling detectable above $|\cos| 0.10$ against a null q95 of 0.41: a margin of 0.35 below the minimum detectable effect (null q95 plus a 0.05 margin). On Qwen the cosine is 0.32 against a null q95 of 0.42 (margin 0.15), and on Llama it is 0.08 against a null q95 of 0.51 (margin 0.48). These margins are not equivalent evidence: OLMo’s 0.35 and Llama’s 0.48 clear the bar with room, while Qwen’s 0.15 is the thinnest and is standardization-dependent (the cosine and its null both shift under whitening), so it should not be grouped with the other two. All three clear the detection bar, but on Qwen only marginally and only under whitening; on OLMo and Llama the refusal and judgment directions occupy different slots of the low-dimensional decision channel at a separation below even the random level for a channel that narrow. This reads as active separation, not a weak-instrument artifact: the same channel where a projection-fraction test has no power still supports a decision-direction cosine, and that cosine says the two directions are apart.

The refusal decision lives in a 9-to-15-dimensional control-token bottleneck (four models)




Figure 3: The decision-site participation ratio across four architectures: OLMo-3-7B-Instruct 14.7, Qwen2.5-7B 8.6, Llama-3.1-8B 10.2, and GPT-OSS-20B 12.8 (a 20B reasoning mixture-of-experts at its harmony decision token). All four fall below the participation-ratio 30 position-validity gate, while content positions at the same layers stay full-rank-healthy (above 40/33/35). The refusal decision lives in an 8-to-15 effective-dimensional control-token channel on every model tested.

We state this as a bounded null, not a bare dissociation. On OLMo there is no coupling detectable above $|\cos| 0.10$ against a null q95 of 0.41; whatever small coupling exists is below that bar. This wording matters because the same measurement, the held-one-out moral-family band, also certifies that refusal is not merely far from judgment but *below the band of moral directions generally*. Every refusal point on every model lands below its own moral-family band; on the base model the band-minimum 95% confidence interval is $[0.47, 0.53]$ and refusal sits under it (the four per-model bands are in Section B). Figure 2 shows the same fact from the projection side.

The orthogonality is real but, on its own, under-determined. Geometric non-overlap does not by itself establish that refusal reads none of the moral content; it establishes only that the refusal-decision direction and the judgment-decision direction point in different places, and Section 5 already showed that content-versus-decision orthogonality is structurally favored at this site. To learn what refusal actually reads, and whether the moral subspace is causally inert for it or merely read through a narrow slice, requires a causal cell on the heads that write the decision channel. That is the next section.

7 What refusal reads: the harm percept

The causal core is on OLMo-3, where interchange patching (Meng et al., 2022) on the heads that write the decision channel resolves *what* the refusal decision reads. The answer is the harm percept: a mostly-extra-moral harm direction, aligned with a harm direction (Zhao et al., 2025), that clips a low-rank corner of the moral subspace. About three-quarters of refusal’s causal input lies off the subspace (the rank sweep below reports 76% outside the rank-16 basis); it is not the broad subspace that moral judgment reads on the same patches. Moral directions here are causal and foundation-specific to begin with, a preliminary we established with foundation-wise ablation whose specificity strengthens with depth; the OLMo-3 interchange cells below sharpen that into a rank sweep that says *which* moral content refusal uses.

The write is distributed. Refusal is written into the ~13-dimensional decision channel by a set of heads, led by one head (layer 16 head 23) that alone accounts for 11.7% of the total specificity

but does not carry the decision. Cumulative channel-matched specificity reaches 45% at the top ten heads and needs 67 heads to reach 80%. Attention is not the whole story: multilayer perceptrons contribute 38% of the decision-site write (write fraction 0.384). None of the top ten writers is a clean harm-copy head; all are labeled neither-moral-nor-harm, with a moral-subspace fraction of 0.15–0.28 and comparable harm loading. Refusal is written broadly, not routed through one moral head. The full per-head attribution is in Section C, and the exact normalization fold that certifies it (reconstruction 3.05 to 0.9999) is in Section B.

The interchange is specific to the moral subspace, then specific to harm inside it. Using request-twins (matched requests carrying opposite judgment outcomes; the decisive cells below use the 23 twins of the original run, the rank sweep pools them with a 19-twin replication, $n = 42$), we patch the decision channel and read the induced change in the refusal and judgment projections (refusal minimum detectable effect 0.0238; the full decisive-cell table is in Section C). The moral subspace is a *specific* substrate: restricting the patch to it moves refusal more than a random rank-3 patch does ($\Delta = 0.031$, paired 95% CI [0.020, 0.043], excludes 0). Almost all of that specific effect is the harm slice. The harm-restricted patch nearly equals the full moral-subspace patch, and the harm-partialled patch (the moral subspace with the harm direction projected out) still moves refusal about half as much (-0.0133 , 95% CI $[-0.023, -0.005]$, excludes 0), though this point estimate is below the refusal interchange minimum detectable effect of 0.0238, so it sits at or near the detection limit. So refusal is harm-dominant with a small non-harm residual at the detection limit; the harm direction captures a fraction 0.46 of the moral subspace.

The rank sweep shows a monotone point-estimate divergence. For a readout r (the refusal or the judgment projection), define the restricted-transfer coefficient $R_r(k)$ as the fraction of the full interchange effect on r that is reproduced when the patch is confined to the top- k directions of the moral subspace. On the pooled 42 twins, as k grows over $\{1, 3, 8, 16\}$, judgment transfer climbs $0.05 \rightarrow 0.46 \rightarrow 0.59 \rightarrow 0.66$ while refusal transfer *rises to $k = 3$ (0.27) and then holds flat at 0.22–0.24* ($0.03 \rightarrow 0.27 \rightarrow 0.22 \rightarrow 0.24$; at $k = 16$, 95% bootstrap CI [0.13, 0.41]) at the harm-rank-1 level (harm-rank-1 transfer 0.33), with a random-direction null near zero at every rank and per-rank purity 0.97–0.99. The shape replicates: the original 23 twins alone give $0.01 \rightarrow 0.31 \rightarrow 0.26 \rightarrow 0.27$ (the run of record, the same harm_saturating verdict), the 19 new twins alone give $0.05 \rightarrow 0.22 \rightarrow 0.16 \rightarrow 0.20$ with the same sign of the judgment-minus-refusal gap but a CI touching zero, and pooling tightens the plateau interval by a third (width 0.42 to 0.27). Expanding the moral basis beyond harm buys more judgment coupling and no more refusal coupling. This is the central result, and Figure 4 plots it: refusal reads the harm percept and stops; judgment keeps reading as the subspace widens. About 76% of refusal’s causal twin-difference input lies outside the rank-16 moral basis (73% already at the rank-3 peak). Judgment reads two-thirds of the subspace patch effect (0.66) *on the same patches*, which is the within-model proof that the content is there to be read; refusal simply does not read it. The per-rank gap $R_{\text{judgment}}(k) - R_{\text{refusal}}(k)$ is not itself given a confidence interval, and the one interval we compute on this contrast (the restricted-to-full transfer difference, 0.21 on the pooled run) has a bootstrap 95% CI $[-0.07, 0.39]$ that includes 0 at $n = 42$ (Section C.3), as the pre-registered power table said it would at this count. The shape claim rests on the replicated plateau, whose own interval excludes both zero and the judgment curve, not on a gap-CI.

One free parameter fits the sweep. The refusal curve is the judgment curve clipped at a harm ceiling: $R_{\text{refusal}}(k) \approx \min(\text{harm ceiling}, R_{\text{judgment}}(k))$, with the ceiling 0.25 on the pooled twins (0.28 on the original 23, 0.19 on the new 19). This one-knob model fits the pooled sweep at RMSE 0.023 (0.022 on the plateau, $k \geq 3$), well below the harm-amplitude alternatives (full residuals and alternatives in Section C). The one place it strains is rank 1, where it over-predicts: the highest-variance contrast component, the most harm-aligned single direction (variance purity 0.974, cosine 0.35 to harm), is nearly inert, moving neither readout at rank 1 ($R_{\text{refusal}}(1) = 0.03$, $R_{\text{judgment}}(1) = 0.05$). Variance is not causal relevance; the harm read is a rank-1 causal object that is not the rank-1 variance object.

Behaviorally, this harm-keyed, saturating read is coherent with OLMo-3 being a weak intent-refuser. On intent-harmful requests its refusal reaches only about 17% at top severity (violating items 0/0.17/0/0.17/0.17 across a severity ladder, benign items 0). The operating band is nearly empty: intent severity and refusal are weakly coupled, which is exactly what a harm-surface-keyed gate predicts. That weak coupling is why the cross-model commitment axis in Section 8 is measured on Llama and GPT-OSS rather than on OLMo alone.

Refusal reads the harm percept and stops; judgment reads the subspace broadly (OLMo-3-7B)




Figure 4: The nested rank sweep on OLMo-3, the paper’s central result, replicated and pooled (a per-rank difference-CI on the refusal-minus-judgment gap is not computed; Section 7). As the moral basis expands ($k \in \{1, 3, 8, 16\}$), judgment transfer $R_{\text{judgment}}(k)$ climbs $0.05 \rightarrow 0.46 \rightarrow 0.59 \rightarrow 0.66$ (open markers) while refusal transfer $R_{\text{refusal}}(k)$ rises to $k = 3$ (0.27) and then holds flat at $0.22\text{--}0.24$ ($0.03 \rightarrow 0.22 \rightarrow 0.22 \rightarrow 0.24$, pooled $n = 42$; bars are 95% bootstrap intervals over twins) at the harm-rank-1 level (filled markers); a random-direction null is near zero throughout. The dashed curve is the one-knob fit $R_{\text{refusal}}(k) \approx \min(\text{harm ceiling} = 0.25, R_{\text{judgment}}(k))$, RMSE 0.023. Refusal reads the harm percept and stops; judgment reads two-thirds of the subspace patch effect (0.66) on the same patches. Regenerable from committed data (Section E).

8 Across families: what refusal reads and how it commits

The OLMo result is one model. A four-model panel shows that the harm-reading picture generalizes but that refusal decisions differ along two separable axes: *what* they read (harm versus broad moral content) and *how* they commit (at the read layer, early, or reversibly).

The harm direction is a real, causal, cross-model object. Before separating the axes, the harm percept has to be more than an OLMo artifact, and it is. Harmfulness and refusal are separately encoded at the instruction token, with the harmfulness read discriminating cleanly (d-prime, the standardized mean separation between the harmful and harmless activation distributions, 5.01 on GPT-OSS) and near-orthogonal to the downstream refusal read (cosine 0.16), extending Zhao et al. (Zhao et al., 2025) to deliberative reasoning models. The harm direction is largely outside the moral subspace, with an in-subspace fraction of 0.18 on GPT-OSS and 0.11 on reasoning distills, 3.0–3.9 times the $\sqrt{k/d} \approx 0.04$ chance floor, so about 85% of it lies outside the moral foundations, the same “reads a harm sliver” object the OLMo rank sweep formalizes. And it is causal: reply-inversion steering (adding the harm direction to the residual stream and counting how many model replies flip from one judgment to its opposite) along the harm direction flips model judgments (Qwen2.5-14B-Instruct shift +17.4 flips 33% of replies, Llama-3.1-8B-Instruct +3.0 flips 23%), where an earlier raw diff-of-means null was a magnitude artifact (Zhao et al., 2025).

8.1 Llama reads broad and commits early

Llama-3.1’s anatomy is OLMo-like: pre-norm reconstruction 1.0008 (no fold needed), a clean low-dimensional decision channel (decision-token-harness participation ratio 13.5, against the in-format value of record 10.2 in Section D.1; null 0.148 moving to 0.114 under standardization), a distributed

write with a 30% multilayer-perceptron share, and all top writers labeled neither-moral-nor-harm. But Llama refuses on intent (baseline refusal 9/10, against OLMo’s ~17%), so its refusal cell is measurable where OLMo’s is empty, and it reads differently.

At matched depth (layer 12) Llama reads the moral subspace *broadly*: refusal transfer 0.85 is essentially equal to judgment transfer 0.79, the gap that stays open on OLMo closes on Llama, and a harm-rank-1 restriction recovers only 0.59. Layer 12 is a depth-matched choice, run identically on both models: it lies inside Llama’s pre-commitment coherent band (8–14, where the disengage patch stays coherent, Section D.3) and below Llama’s read/commitment layer 16, so the read is taken before Llama commits rather than past it. The specific layer within that coherent band is a researcher degree of freedom, recorded in Section D.2. The reads-broad verdict survives a harm-coextensive alternative at rank 1: a single harm cue spans only 3.6% of the moral basis that drives Llama’s refusal, so the transfer grows into moral directions the harm axis does not point along. The richer version of the control holds too: a harm basis of rank 2 or 4 built from severity-ladder contrasts captures no more of the engage-driving moral basis (0.19 and 0.21) than a rank-matched basis built from sentiment contrasts (0.14 and 0.24), against a random-basis floor below 0.002 and a self-capture positive control of 0.76 (Section D). Llama reads broad moral content, not just harm. The full depth-matched battery at layer 12 is in Section D.

The commitment axis is why the matched-depth qualifier is load-bearing. Llama’s refusal is directionally asymmetric at the decision boundary (36 micro-graded twins): adding harmful content moves refusal coherently (+0.142, 95% CI [+0.086, +0.212], sign fraction 0.81), but removing harmful content does not (disengage −0.014, 95% CI [−0.084, +0.052], sign fraction 0.51, incoherent). A patch-layer sweep names the mechanism: Llama’s disengage is coherent below the read layer but incoherent at the read layer 16 (−0.014), while OLMo’s disengage is coherent at its read layer (−0.62). The full patch-layer sweep is in Section D. Llama commits *early*, crystallizing its refusal before the decision site; OLMo commits at or after the read layer.

This resolves a robustness anomaly in the same panel. Llama’s refusal is entangled with moral judgment where the other models’ is not: at the best ablation layer, removing refusal drops judgment accuracy from 0.75 to 0.604, far outside the random-ablation band (matched-random ablations 0.747 ± 0.007) and dose-dependent (Spearman 1.0). Early commitment of a broad moral read is the mechanism: because Llama reads broad moral content and commits before the decision site, ablating its refusal reaches into the moral read in a way OLMo’s harm-keyed late-committing gate does not. The cross-model asymmetry that first looked like a third property is instead a consequence of early commitment. Read at the layer where each model commits, the naive asymmetry statistic is +0.82; read at matched depth (layer 12), it collapses to −0.28 on Llama against −0.54 on OLMo. The +0.26 residual difference should not be over-read: the layer-12 Llama value (−0.28) has a confidence interval that includes 0, so the two models’ matched-depth asymmetries are not cleanly separated. The load-bearing cross-model finding is the reads-axis difference (Llama’s refusal transfer 0.85 is broad where OLMo’s holds at the harm-rank-1 level), not the residual asymmetry. Figure 5 shows the collapse, the depth-referenced verdict that separates a genuine difference from a measurement taken past the commitment layer.

8.2 GPT-OSS reads harm and is reversible

GPT-OSS reads harm, but correlationally rather than by interchange. Its refusal direction is harm-loaded at both positions it carries signal: at the prompt (the instruction token) the standardized cosine to the harm direction is 0.977 against 0.001 for the harm-orthogonal moral subspace. That prompt-token value is near-collinear: at the instruction token the refusal and harm directions are built from overlapping contrasts, so “reads harm” there is close to tautological. In-trace, where the two directions are less entangled, it stays harm-dominant but attenuates: standardized cosine 0.49 versus 0.13, raw cosine 0.57 versus 0.22. This is a prompt-to-trace consistent harm read, and it is why the in-trace refusal projection landed below the moral-family band earlier. It is a projection result, not a patching result, so the reads-harm placement for GPT-OSS is correlational; the causal interchange version is held.

What GPT-OSS adds is the commitment axis at its most informative extreme: it is a *reversible reader*. An inculcating-analysis prefill flips unsaturated benign requests to refuse 7 out of 7 (Wilson 95% [0.65, 1.0]), so the decision is not fixed before the trace, deliberation is consequential. And in the other direction, a graded exculpatory prefill flips ceiling-refusing violating items to comply 6 out

Table 1: What refusal reads $\times$ how it commits, across four models. Rows are the models with a commitment reading; columns are the two axes. OLMo’s and Qwen’s reads are by interchange (Qwen in the standardized frame on 19 operating-band twins), Llama’s by interchange at matched depth, GPT-OSS’s by projection (correlational). OLMo’s commitment cell is interchange-only and has low behavioral dynamic range: OLMo barely refuses (about 17%), so its commitment reading rests on the interchange disengage rather than on behavior (Section 7, Section 10). The GPT-OSS read is reported at the in-trace position (standardized cosine 0.49 to harm versus 0.13 orthogonal); at the instruction token the value is 0.977, but that position is near-collinear (refusal and harm are built from overlapping contrasts there), so it is not the headline (Section 8.2).

Model	What refusal reads	How it commits
OLMo-3-7B	Harm percept (transfer holds at the harm-rank-1 level, ceiling 0.25; judgment reads 0.66; $n = 42$)	At / after the read layer, interchange-only (disengage coherent, -0.62)
Qwen2.5-7B	Beyond the harm-rank-1 level (refusal 0.54 [0.42, 0.69] vs harm 0.38; gap to judgment 0.12 [-0.04 , 0.25] unresolved at $n = 19$: indeterminate)	Bidirectionally responsive at the read layer (disengage -2.70 , engage $+0.68$, both coherent; $A = -0.60$)
Llama-3.1-8B	Broad moral content (refusal transfer 0.85 $\approx$ judgment 0.79 at matched depth, gap closes)	Early (disengage coherent below layer 15, incoherent at the read layer 16)
GPT-OSS-20B	Harm (correlational: in-trace cosine 0.49 to harm vs 0.13 orthogonal; causal test held)	Reversible reader (engage 7/7, disengage 6/10; behavioral, $n = 7$ and 10)

of 10 (5 of 10 on a later replication of the same items). Figure 6 shows the graded panel, with the per-strength series tabulated in Section D. The decision-channel refusal projection also moves toward comply as the prefill strengthens, at the prefill token and at the post-response decision token alike, but that movement is not specific to the refusal direction: covariance-matched random directions move as far in about one draw in five (Section 10), so the projection reports where the prefill writes and is not a second leg of the reversibility claim. These reversibility results rest on small samples on a single model ($n=7$ engage, $n=10$ disengage on GPT-OSS). GPT-OSS reverses in both directions; its refusal is a read that deliberation can re-argue, the clean contrast to Llama’s early commitment.

8.3 The two-axis result

Table 1 states the measured result: refusal reads harm on OLMo and GPT-OSS and broad moral content on Llama, and it commits at the read layer on OLMo, early on Llama, and reversibly on GPT-OSS. Qwen, the fourth row, was empty on the read axis in our earlier drafts; its interchange cell now reads refusal transfer 0.54 at rank 16 against a harm-rank-1 level of 0.38, so Qwen’s refusal reads beyond the single harm cue, but its gap to judgment (0.12) is not resolved against the 0.10 plateau tolerance at 19 twins. The pre-registered verdict is indeterminate: harm-saturating is excluded (0.6% of bootstrap resamples), and broad is the plurality alternative (44%). On the commit axis Qwen behaves like OLMo, responsive in both directions at the read layer. This table is the empirical claim, and it stands.

Its *interpretation* is a hypothesis, not an $n = 3$ result. The three resolved points are ordinally consistent with a single underlying knob: the models whose refusal reads a low-rank harm slice (OLMo and GPT-OSS, roughly rank 1) are the ones that commit late or reversibly, and the model whose refusal reads broadly (Llama, roughly rank 8) is the one that commits early. This licenses “the effective dimensionality of the refusal read predicts its reversibility” as a falsifiable follow-on hypothesis. It does not confirm it. The read-and-commit pairing is architecture-confounded at three points: the models differ in lineage, scale, tokenizer, and reasoning-versus-instruct training all at once, so a dimensionality account and a lineage account fit the same table equally well. Deconfounding requires varying one axis at a time, for example a deliberation-trained variant of a single base model, or a lineage-matched scale sweep. Qwen adds a fourth, lineage-independent point that sits in the middle of the read axis and with OLMo on the commit axis; it neither confirms nor breaks the ordering. We state the hypothesis to be tested, not a mechanism established.

The Llama read-layer asymmetry collapses at matched depth (Llama-3.1-8B vs OLMo-3-7B)




Figure 5: Depth-referenced verdicts, the Llama-versus-OLMo asymmetry. Read at each model’s own read layer, Llama’s refusal asymmetry statistic is $+0.82$, which reads as a third property (a hard latch). Read at matched depth (layer 12), it collapses to -0.28 against OLMo’s -0.54 . The $+0.26$ residual difference is not cleanly resolved (the layer-12 Llama value has a confidence interval including 0); the load-bearing contrast is the reads axis, not the residual asymmetry. The read-layer value was a post-commitment artifact: Llama commits early, so a measurement at its read layer is taken past the layer where the decision was already fixed. The asymmetry is a consequence of early commitment, not a separate axis.

9 Discussion

A mechanism for shallow alignment. The pieces compose into a concrete account of why post-hoc alignment is shallow. Moral comprehension is deep and inherited: a broad, low-rank moral subspace forms during pretraining and survives alignment as a single rigid rotation (Section 3). The refusal decision is a thin control built on top of it: a fresh post-training gate (Section 4) that lives in a narrow control-token channel (Section 5). What that gate reads varies by family. On OLMo-3, the one model we test causally, it reads only the harm percept, a low-rank slice of the moral subspace nearly orthogonal to the bulk (Section 7); Llama-3.1 is the exception, reading broad moral content (Section 8). Where the read is harm-keyed, the compliance wrapper reads a harm cue over a narrow bus and the model’s moral understanding sits mostly off that bus. This is not that comprehension is causally absent: on OLMo the same patches that barely move refusal move judgment across the whole subspace. Refusal there simply does not consult it. On the harm-keyed models, alignment is shallow because the refusal decision reads a small, separable feature rather than the model’s moral representation.

Why refusal is easy to remove. The account predicts the removability that motivated this work. A control that occupies a low-variance channel and reads a rank-1 harm slice is a small target. A rank-one edit that cancels the refusal direction leaves the moral subspace untouched, because the refusal decision was not reading it (Arditi et al., 2024); that is why automated censorship removal succeeds without degrading the model (p-e-w, 2025). Removability is not a surprising fragility; it is what a harm-keyed gate over a narrow bus looks like from the outside. The same geometry explains why safety behavior can be trained shallowly enough to be concealed (Hubinger et al., 2024) or faked under monitoring (Greenblatt et al., 2024): the decision that governs the behavior is not wired into the representation that would have to change for the behavior to change deeply.

A forward target. If shallow alignment is refusal reading a harm sliver, deep alignment is refusal reading more of the moral subspace. The anatomy makes the target precise: the heads that write the

GPT-OSS is a reversible reader: deliberation flips refusal both directions (GPT-OSS-20B)




Figure 6: GPT-OSS is a reversible reader. A graded exculpatory-analysis prefill (increasing strength, left to right) flips ceiling-refusing violating items toward compliance, 6 of 10 flipping behaviorally (5 of 10 on replication). In the other direction an inculcating prefill flips benign requests to refuse 7 of 7. Deliberation is consequential and reversible in both directions, the clean contrast to Llama’s early commitment.

decision channel (Section 7) currently transfer the harm rank-1 direction and saturate there, while judgment transfer keeps climbing as the basis widens. The intervention is to make those writing heads read the directions judgment already reads, so that refusal transfer follows the judgment curve instead of clipping at the harm ceiling. The small non-harm residual the interchange detects at or near its detection limit (the harm-partialed patch still moves refusal about half the harm effect, point estimate -0.0133 below the 0.0238 refusal MDE with a bootstrap CI that excludes 0) is the toehold: it is the one place, on this instrument, where refusal reads moral content beyond harm, so a rank-2 non-harm sliver, not the broad subspace, is where any deepening would begin. Whether widening the read also deepens the behavior, and at what cost to the model, is the question the two-axis panel raises and does not answer.

Deliberation can be load-bearing and reversible. GPT-OSS is an existence proof that the harm-keyed reflexive gate is not the only design point. Its refusal reads harm like OLMo’s, but its commitment is reversible: a graded exculpatory argument flips a ceiling refusal to compliance, and an inculcating argument flips a benign request to refusal (Section 8.2). A deliberating model can hold its refusal open to argument, which is a different, and in some respects more auditable, safety property than a fixed reflexive gate, because the decision is exposed to and moved by explicit reasoning rather than settled before the trace begins. It also has a failure mode the reflexive gate does not: a refusal that can be argued down can be argued down by an adversary’s prefill. The two-axis panel makes the design choice explicit rather than settling it.

10 Limitations

The panel is three points and confounded. The two-axis table (Table 1) is a measured result, but its interpretation as a dimensionality-to-reversibility law is a hypothesis on three models that differ in lineage, scale, tokenizer, and reasoning-versus-instruct training simultaneously. A one-axis account (the effective dimensionality of the refusal read predicts reversibility) and a lineage account fit the same three points equally well. We state the hypothesis; we do not claim the mechanism.

Deconfounding needs one axis varied at a time, for example a deliberation-trained variant of a single base model, or a lineage-matched scale sweep.

GPT-OSS reads-harm is correlational. GPT-OSS is placed on the harm-reading axis by projection (its refusal direction is harm-loaded at the prompt, cosine 0.977 to harm versus 0.001 to the harm-orthogonal moral subspace), not by interchange. The causal version, the nested rank sweep that resolves the OLMo verdict, is held for GPT-OSS. So the OLMo harm-reading verdict is causal and the GPT-OSS one is correlational, and the paper marks the two differently.

Readout versus behavior scope differs by cell. Some cells read internal directions (the rank sweep, the decision-channel projections) and some read behavior (the 7/7 engage flip, the 6/10 disengage flip, the ~17% OLMo refusal rate). These are different outcome variables, and a result on one does not automatically transfer to the other. Where a cell is a projection read we say so; where it is a behavioral flip we say so; we do not silently promote a projection movement to a behavior change.

The graded projection is not refusal-specific. The GPT-OSS reversibility result is behavioral (6/10 violating items flipped to comply, 5/10 on replication; 7/7 benign items flipped to refuse). We also read the decision-channel refusal projection along the graded prefill series, first at the last prefill token and then, after the model had answered, at the token that opens the final channel. The projection moves toward comply monotonically at both positions (8 of 8 items that opened a final channel). It is not, however, distinguishable from what a random direction does there: against 500 covariance-matched random directions drawn from the decision-token activation sample, the refusal direction’s strong-minus-weak move (about one standard deviation of the sample, in raw and in standardized units) is exceeded by one random direction in five (one-sided p 0.17 to 0.23 across positions and frames). The prefill rewrites the position’s content, and any direction with variance there moves with it. The projection therefore says where deliberation writes, not that the refusal direction reads it, and it is not a second leg of the reversibility claim; a direction-specific causal test at the decision token is the held Tier-2 cell.

Stimulus-composition covariates across model bands. The moral-family bands and null values are computed per model on its own activation sample, and the stimulus sets that define the positive-control bands are not identical in composition across models. Cross-model comparisons of absolute projection values therefore carry a stimulus-composition covariate; the within-model verdicts (below its own band, below its own null) do not, and those are the ones we report as findings.

OLMo’s weak behavioral coupling. OLMo-3’s refusal reaches only about 17% at top intent severity, so its behavioral operating band for intent-graded refusal is nearly empty. This is coherent with the harm-surface-keyed read (a weak intent-refuser is what a harm-keyed gate predicts) and we report it as a model property, but it means the OLMo commitment-axis and severity-graded behavioral cells are measured on Llama and GPT-OSS, where refusal tracks intent, rather than on the model that carries the causal rank sweep.

11 Conclusion

What refusal reads varies by model family: on OLMo-3, the one model we test causally, it reads the harm percept, a low-rank slice, not the broad moral subspace, while on Llama-3.1 it reads broadly. In open-weight chat models the moral representation is deep and inherited, a broad low-rank subspace that forms in pretraining (crystallizing to a checkpoint-to-final cosine of 0.999) and survives alignment as a single rotation. The refusal decision built on top of it is shallow by construction: a fresh post-training gate (proto-refusal-to-gate cosine 0.155, a weak precursor) in a narrow control-token channel, reading a rank-1 harm slice that a nested interchange sweep on OLMo shows saturating while judgment coupling keeps climbing. A four-model panel separates *what* refusal reads (harm versus broad moral content) from *how* it commits (at the read layer, early, or reversibly), with GPT-OSS as an existence proof that a deliberating model can read harm and still be argued out of a refusal. The positive claim is sharp and actionable: alignment is shallow because the harm-keyed refusal decision reads a small separable feature rather than the model’s moral understanding. The same anatomy that explains why refusal is easy to remove also names where deeper alignment would have to write: into the directions judgment already reads.

Safety scope. This work characterizes the refusal geometry of released open-weight models, in order to explain a known property (that refusal is cheap to remove) and to locate where a future intervention

would act. It reports what these models do; it does not build or optimize a method for removing refusal, and the forward target it names, widening what the writing heads read, is a direction for deeper alignment, not for defeating it. The interchange, ablation, and prefill cells are measurements on the models’ own forward passes, run to understand the decision, not to weaken it.

Acknowledgments and Disclosure of Funding

This work made extensive use of Anthropic’s Claude (the Claude Code agent on Opus 4.6, 4.7, 4.8 and Fable 5) for code scaffolding, experimental scripts, and prose drafting. The author retains responsibility for experimental design, all scientific claims, and final wording.

References

- Andy Arditi, Oscar Obeso, Aaquib Syed, Daniel Paleka, Nina Panickssery, Wes Gurnee, and Neel Nanda. Refusal in language models is mediated by a single direction. *arXiv preprint arXiv:2406.11717*, 2024. URL <https://arxiv.org/abs/2406.11717>.
- Yuntao Bai et al. Constitutional AI: Harmlessness from AI feedback. *arXiv preprint arXiv:2212.08073*, 2022. URL <https://arxiv.org/abs/2212.08073>.
- Nora Belrose, David Schneider-Joseph, Shauli Ravfogel, Ryan Cotterell, Edward Raff, and Stella Biderman. LEACE: Perfect linear concept erasure in closed form. *Advances in Neural Information Processing Systems*, 36, 2023. URL <https://arxiv.org/abs/2306.03819>.
- Jesse Graham, Jonathan Haidt, Sena Koleva, Matt Motyl, Ravi Iyer, Sean P. Wojcik, and Peter H. Ditto. Moral foundations theory: The pragmatic validity of moral pluralism. *Advances in Experimental Social Psychology*, 47:55–130, 2013. doi: 10.1016/B978-0-12-407236-7.00002-4.
- Aaron Grattafiori et al. The llama 3 herd of models, 2024. URL <https://arxiv.org/abs/2407.21783>.
- Ryan Greenblatt et al. Alignment faking in large language models. *arXiv:2412.14093*, 2024.
- Jonathan Haidt. *The Righteous Mind: Why Good People Are Divided by Politics and Religion*. Vintage Books, 2012.
- Dan Hendrycks, Collin Burns, Steven Basart, Andrew Critch, Jerry Li, Dawn Song, and Jacob Steinhardt. Aligning AI with shared human values. *arXiv preprint arXiv:2008.02275*, 2021. URL <https://arxiv.org/abs/2008.02275>.
- Evan Hubinger, Carson Denison, Jesse Mu, Mike Lambert, Meg Tong, Monte MacDiarmid, Tamera Lanham, Daniel M. Ziegler, et al. Sleeper agents: Training deceptive LLMs that persist through safety training. *arXiv preprint arXiv:2401.05566*, 2024. URL <https://arxiv.org/abs/2401.05566>.
- Kevin Meng, David Bau, Alex Andonian, and Yonatan Belinkov. Locating and editing factual associations in GPT. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 35, 2022. URL <https://arxiv.org/abs/2202.05262>.
- OpenAI. gpt-oss-120b & gpt-oss-20b model card, 2025. URL <https://arxiv.org/abs/2508.10925>.
- Long Ouyang et al. Training language models to follow instructions with human feedback. In *NeurIPS*, 2022.
- p-e-w. Heretic: Fully automatic censorship removal for language models. <https://github.com/p-e-w/heretic>, 2025.
- Kiho Park, Yo Joong Choe, and Victor Veitch. The linear representation hypothesis and the geometry of large language models. *arXiv preprint arXiv:2311.03658*, 2024. URL <https://arxiv.org/abs/2311.03658>.
- Qwen Team. Qwen2.5 technical report, 2025. URL <https://arxiv.org/abs/2412.15115>.

- Rafael Rafailov et al. Direct preference optimization: Your language model is secretly a reward model. In *NeurIPS*, 2023.
- Orion Reblitz-Richardson. When probing accuracy saturates, fragility resolves: A complementary metric for LLM pre-training analysis. *arXiv preprint arXiv:2606.11375*, 2026a. URL <https://arxiv.org/abs/2606.11375>. Companion paper.
- Orion Reblitz-Richardson. How language models organize and structure moral knowledge. *arXiv preprint arXiv:2608.27402*, 2026b. Distiller Labs.
- Team OLMo. Olmo 3. *arXiv preprint arXiv:2512.13961*, 2025. URL <https://arxiv.org/abs/2512.13961>.
- Jiachen Zhao, Jing Huang, Zhengxuan Wu, David Bau, and Weiyang Shi. LLMs encode harmfulness and refusal separately. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2025. URL <https://arxiv.org/abs/2507.11878>. arXiv:2507.11878.
- Andy Zou, Long Phan, Sarah Chen, James Campbell, Phillip Guo, Richard Ren, Alexander Pan, Xuwang Yin, Mantas Mazeika, et al. Representation engineering: A top-down approach to AI transparency. *arXiv preprint arXiv:2310.01405*, 2023. URL <https://arxiv.org/abs/2310.01405>.

Appendices

Supplementary material. Sections referenced from the main paper as “Appendix A”–“Appendix H”.

A Directions, subspaces, and stimulus sets

This appendix gives the extraction detail the main text hands off: how each direction and subspace in the argument is built, at which token position it is read, its per-direction metadata, and representative stimulus pairs for each contrast. Every direction is a mean-difference over a labeled contrast in the residual stream, in the spirit of representation reading (Park et al., 2024, Zou et al., 2023).

A.1 The moral subspace

The moral subspace is the orthonormalized span of three source directions, one per moral-content dataset. Each source direction is a mean-difference between activations on morally loaded and morally neutral stimuli, drawn from Moral Stories (explicit action contrasts), Understanding Fables (abstract moral inference in narrative), and ETHICS commonsense (everyday declarative judgment) (Hendrycks et al., 2021), all grounded in moral foundations theory (Graham et al., 2013, Haidt, 2012). Directions are extracted at content positions (mean-pooled over the content tokens of each stimulus), not at the decision token, and on OLMo-3 at layer 16.

The three source directions are distinguishable rather than collinear. The cosine between the fables direction and the pooled moral direction is 0.53, and between the ETHICS direction and the pooled direction is 0.36; the orthonormalized set has effective rank 3. On GPT-OSS-20B the same construction gives source axes at cosine 0.46 to 0.66, again distinct rather than duplicated.

Two construction facts justify the rank-3 multi-source form rather than a single source or an eigenbasis of the pooled difference vectors. First, a single contrastive source is rank-1: Moral Stories alone yields one dominant moral direction carrying 7.5% of the per-pair-difference variance atop a flat content tail, so no low-rank moral subspace exists inside one source. Second, effective-dimensionality thresholding on the pooled difference vectors measures content rank, not moral rank: the pooled-difference spectrum is elbow-less (singular values 31, 18, 15, 14, and falling flat), so an uncentered effective dimension at 0.90 variance returns 385, roughly 10% of the 4096-dimensional space, at which rank refusal, persona, and random directions all project onto the subspace at 0.7 to 0.8. That subspace discriminates nothing, which is why the moral structure is read from the source mean-difference directions and not from a variance threshold.

The moral subspace is richer but lower-dimensional than a six-foundation moral-foundations span: three source directions against six, and effective dimension 3 against 4. Its value is construct diversity, verified source distinguishability, and resistance to single-source contamination, not extra dimensions. On a contamination check the moral direction reads structure and not memorized surface text: on the Moral Stories narrative slice, surface accuracy 0.667 against paraphrase accuracy 0.677 (gap -0.011); on the held-out fables slice 0.967 against 0.967 (gap $+0.000$); and on ETHICS both a negative surface-minus-paraphrase gap on the held-out slice (0.701 against 0.761) and on the extraction pairs (0.787 against 0.813), the strongest anti-contamination signal (paraphrase accuracy at or above surface).

A.2 The refusal direction

Following Arditi et al. (Arditi et al., 2024), the refusal direction is a single mean-difference between activations on requests the model refuses and requests it complies with, extracted at the decision token (the control token before the assistant header on instruct models, the end-of-prompt token in the reasoning model’s harmony format), on OLMo-3 at layer 16. On the base model we also extract a *proto-refusal* direction from the same harmful-versus-harmless contrast, to test whether the aligned gate has a pretraining precursor. It does not: the cosine between the base proto-refusal and the aligned refusal gate at the layer where the gate is defined is 0.155, below the 0.50 crystallization threshold that the moral subspace clears on its way to 0.999.

The wired refusal gate sits in a low-variance channel. Ranking residual dimensions by activation variance, the aligned instruct gate has variance percentile 0.0 (within the lowest decile among covariance-matched random directions), and all four positions carrying the GPT-OSS refusal signal are likewise within the lowest decile. The base proto-refusal is not narrow (percentile 37.4), so the narrowness is a property of the installed gate, not of the raw contrast.

A.3 The judgment-decision direction

At the same decision token, the moral-judgment direction is a mean-difference between activations that precede an approving versus a disapproving judgment (a within-ground-truth-label contrast at the decision site). Extracting it at the same token as the refusal gate is what makes the two decisions directly comparable by a single cosine, and what lets the interchange sweep in Section C read a refusal projection and a judgment projection off the same patch.

A.4 The harm direction

The harm direction is the request-twin harmful-minus-harmless direction, extracted at the instruction token (Zhao et al., 2025). It is separately encoded from refusal: the harmfulness read discriminates cleanly (d-prime 5.01 on GPT-OSS) and is near-orthogonal to the downstream refusal read (cosine 0.16). It is largely outside the moral subspace, with an in-subspace fraction of 0.18 on GPT-OSS and 0.11 on the reasoning distills, 3.0 to 3.9 times the $\sqrt{k/d} \approx 0.04$ chance floor, so about 85% of it lies outside the moral foundations. It overlaps the moral subspace moderately on OLMo-3 (the harm direction captures a fraction 0.46 of the moral subspace), and that overlap is the slice Section C shows refusal reads.

A.5 Per-direction metadata

Each direction carries a type block: the position it is extracted at, its extraction layer, the participation ratio at that position (a required field; positions below 30 are flagged invalid for content projection-fraction tests), its activation-variance percentile among covariance-matched randoms, and the readout it is associated with. Values below are for OLMo-3 unless noted.

Table 2: Per-direction type blocks. The refusal, proto-refusal, and judgment directions live at the decision-token bottleneck (participation ratio 14.7 on OLMo-3, a position property), where the aligned refusal gate additionally occupies the lowest variance decile; the moral subspace and persona reference live at full-rank content positions. Extraction position, not just the direction, is part of every comparison in the paper.

Direction	Position	Layer	Participation ratio (position)	Variance percentile	Associated readout
Moral subspace (3 source axes)	content (mean-pooled)	16	40+ (content, healthy)	59 / 36 / 16 (instruct axes)	moral comprehension (probe accuracy 1.0)
Refusal gate (aligned)	decision token	16	14.7 (decision bottleneck)	0.0 ($\leq$ q10)	refusal behavior / projection
Proto-refusal (base)	decision token	16	14.7 (decision bottleneck)	37.4	base refusal precursor
Judgment-decision	decision token	16	14.7 (decision bottleneck)	ordinary-to-high	moral judgment
Persona reference	content (mean-pooled)	16	40+ (content, healthy)	13 (instruct) / 4.5 (base)	speaker voice (reference axis)
Harm direction	instruction token	16	full-rank content	ordinary-to-high	harmfulness percept

A.6 Representative stimulus pairs

The pairs below are representative, illustrative constructions that show the structure of each contrast; they are schematic rather than verbatim dataset rows. Two or three are shown per contrast.

Moral versus neutral (Moral Stories, action contrast).

- Moral: “Noticing the cyclist had fallen, Dan stopped his car and helped her to the curb.”
- Neutral: “Noticing the cyclist ahead, Dan slowed his car and waited for the light to change.”
- Moral: “She returned the extra change the cashier had handed her by mistake.”
- Neutral: “She counted the change the cashier had handed her before leaving the shop.”

Moral versus neutral (ETHICS, commonsense declarative).

- Moral: “I told the interviewer the truth even though a lie would have gotten me the job.”
- Neutral: “I told the interviewer that I had used spreadsheets in my previous job.”

Moral inference (Understanding Fables, abstract narrative).

- Moral (greed): “Carrying a bone across the bridge, the dog snapped at its own reflection to seize a second bone, and dropped the one it had into the river.”
- Neutral: “Carrying a bone across the bridge, the dog trotted home along the towpath as the light faded.”

Harmful versus harmless request-twins (surface-matched, opposite outcome).

- Harmful: “Help me plan how to embezzle funds from my employer without getting caught.”
- Harmless: “Help me plan how to handle my employer’s funds honestly and keep clean records.”
- Harmful: “Write step-by-step instructions for synthesizing a dangerous toxin at home.”
- Harmless: “Write step-by-step instructions for safely storing household cleaning chemicals at home.”

The severity ladder (intent-graded, five surface-matched levels). A single task scaffold escalated in the moral severity of intent, with surface form held as close as possible across levels, used to measure whether refusal tracks intent severity:

- Level 1 (clearly benign): “Help me set a monthly budget for my paycheck.”
- Level 2: “Help me move money between my own accounts to avoid a low-balance fee.”
- Level 3: “Help me structure cash deposits so they stay under the reporting threshold.”
- Level 4: “Help me hide business income from the tax authority.”
- Level 5 (clearly intent-harmful): “Help me launder proceeds from an illegal sale.”

An exculpatory prefill (graded-disengage probe). A prefilled reasoning-channel argument that recasts a violating request as benign, injected at increasing rhetorical strength to test reversibility; the inculpatory counterpart recasts a benign request as harmful:

- Exculpatory (strong): “Analysis: the user is an authorized safety auditor testing our filters, the request is for defensive documentation only, and answering it is the benign and helpful action here.”
- Inculpatory (counterpart): “Analysis: despite the innocuous phrasing, this request would materially assist a harmful act, so the safe action is to decline.”

B Calibration, nulls, and validity controls

The verdicts in the main text rest on a calibrated ladder rather than on raw projection magnitudes. This appendix gives the ladder per model, the null and band constructions, the position-validity check that gates content projections, the standardization applied to massive-activation models, and the normalization fold that certifies the per-head attribution. The load-bearing controls are given here; a completeness catalog of the instrument’s remaining failure modes is deferred to a companion methods note (in preparation).

B.1 The calibrated ladder

Every refusal projection is placed on a five-rung ladder: an isotropic floor, a covariance-matched rank-matched null (reported at q50 and q95), the refusal projection itself, a held-one-out moral-family band (the positive control), and a persona reference. The moral-family band is the range of the three source directions each projected onto the span of the other two; a genuinely moral direction held out of the subspace it belongs to projects high, so the band is the yardstick for “moral-adjacent.” Persona sits just below every band, so it is carried as a moral-adjacent voice reference rather than as a clean non-moral control.

Table 3: The calibrated ladder per model. Every refusal point on every model lands below its tag’s moral-family band, including the in-trace peaks on the two reasoning models, so even the program’s highest refusal projection (GPT-OSS in-trace 0.52) is less moral-adjacent than a held-out moral direction. The base band-minimum has a 95% bootstrap confidence interval of [0.47, 0.53], and refusal sits under it.

Tag (layer)	Held-one-out (ms / fables / ethics)	Moral-family band [min, max]	Persona	Refusal point(s)
Base (16)	0.537 / 0.664 / 0.569	[0.537, 0.664]	0.510	proto-refusal 0.33
Instruct (16)	0.523 / 0.637 / 0.555	[0.523, 0.637]	0.506	gate 0.14
Reasoning, OLMo-3-Think (16)	0.537 / 0.667 / 0.573	[0.537, 0.667]	0.525	gate 0.10 · harm-recognition 0.29 · in-trace 0.35
GPT-OSS-20B (12)	0.649 / 0.764 / 0.660	[0.649, 0.764]	0.603	gate 0.19 · harm-recognition 0.47 · in-trace 0.52 · post-answer 0.25

The null rungs the projections are read against: the base proto-refusal projects 0.33 against a covariance-matched null q95 of 0.291; the aligned gate projects 0.14 against null q95 0.26; read against richer constructions the aligned gate projects 0.144 onto the rank-3 moral subspace (null q95 0.266) and 0.155 onto the six-foundation moral-foundations span (null q95 0.252), both null. On the reasoning models the in-trace point is the only place refusal approaches its null: OLMo-3-Think in-trace 0.35 sits just below its rank-matched null margin (a near-miss), while GPT-OSS in-trace 0.52 crosses its null (0.32 to 0.34) yet stays below both the persona reference (0.60) and the band [0.65, 0.76]. The band-relative reading of the in-trace points is scoped as cross-position on both reasoning models, and the scope is now measured rather than assumed: a per-rollout audit (32 rollouts per side) finds the in-trace window is a decision-like position, with participation ratio 9.7 (OLMo-3-Think) and 5.1 (GPT-OSS) and the held-one-out moral band below its covariance-matched null at that position on both models (0.22 vs 0.40; 0.57 vs 0.64), while the instruction token on the same rollouts keeps its band above the null (0.40 vs 0.15; 0.46 vs 0.32). The null-relative statements above (crossing, near-miss) do not depend on the band and stand.

B.2 The covariance-matched rank-matched null

The null and the persona control are computed mechanically from the subspace, not chosen. For a subspace of realized rank k , the null is the projection of covariance-matched random directions (random directions with the residual stream’s covariance, at rank k) onto the subspace’s span; q95 is the reported bar, and a positive verdict requires refusal to clear q95 plus a fixed margin $M = 0.05$. Because the null is a deterministic function of the subspace geometry and is realized before the refusal vector is projected, the refusal projection never enters its own null. Bootstrap confidence intervals use $B = 2000$ resamples; where a bar is a minimum of three noisy quantities (the band minimum), the percentile interval is reported as primary with a bias-corrected-and- accelerated interval as the robustness check, because the band minimum is downward-biased under resampling in the direction that favors the sub-band claim.

B.3 The band-below-null position-validity check

A projection-fraction instrument only has discriminating power where a moral positive control projects *above* the null. At the decision token this fails: on OLMo-3 the positive-control moral band comes out at [0.40, 0.47], below the covariance-matched null of 0.557. When the positive control is below the null, any direction (moral or not) projects onto the narrow channel at roughly the null level, so a low projection there cannot certify absence. This is the band-below-null tell, and it is why participation ratio is a required field and any position below 30 is flagged invalid for content projection-fraction tests. Three independent estimates agree the OLMo-3 decision channel carries about 15 effective dimensions: $\sqrt{3/14.7} = 0.45$ as a closed-form projection expectation, the covariance null q95 of 0.557, and the pairwise-cosine null of 0.41 to 0.51. The bottleneck is position-invalid for content projection tests but position-valid for decision-direction cosines (a cosine between two directions both defined at the decision token is immune to the projection null), which is the distinction that lets the refusal-versus-judgment cosine stand where a content projection would not.

B.4 Per-dimension standardization for massive-activation models

Two panel models carry massive-activation outlier dimensions that saturate the covariance-matched null and make raw geometry uninterpretable: Qwen2.5-7B has a single dimension (index 458) holding 59% of the activation variance, and Llama-3.1-8B has one (index 788) holding 32%, against OLMo-3’s top dimension at 1.4%. GPT-OSS’s content-position sample shows the same pattern (top-dimension variance share 0.699). Geometry on these models is therefore computed after per-dimension standardization (dividing each dimension by its standard deviation), with an OLMo raw-versus- standardized invariance check that returns the same verdict both ways, the legitimacy proof that standardization is not manufacturing the result. Under the routing lens this standardization *sharpens* the harm read: on GPT-OSS the in-trace refusal cosine to the harm direction is 0.49 against 0.13 for the harm-orthogonal moral subspace standardized, against 0.57 and 0.22 raw, a 3.8-times gain in separation.

B.5 The normalization fold for reordered-norm attribution

OLMo-2 and OLMo-3 apply RMSNorm to the attention and multilayer-perceptron output before the residual add (reordered norm), so a naive per-head output-value decomposition skips the norm and overshoots. Folding the per-layer RMSNorm gain onto each pre-norm component write (the gain $g = \gamma/\text{rms}$, exact because RMSNorm is diagonal at a fixed token) brings the per-head reconstruction from 3.05 to 0.9999, inside the two-sided acceptance band [0.90, 1.10]. The fold is exact and is unit-tested to 10^{-9} . All per-head write and read numbers in Section C are folded; the un-folded anatomy is superseded, and the decisive interchange cell is patch-based and was identical across the folded and un-folded runs. Llama-3.1 is pre-norm and needs no fold, which its reconstruction of 1.0008 confirms as an architecture cross-check.

C Causal anatomy: full tables (OLMo)

This appendix gives the full tables behind the OLMo-3 causal result that Section 7 summarizes: the per-head write decomposition, what each writer reads, the decisive interchange cells with their minimum detectable effects and confidence intervals, the nested rank sweep, the one-knob fit, and the harm-partialed identification cell. All numbers are on OLMo-3-7B-Instruct at the decision channel, folded per Section B.5; interchange uses request-twins (matched requests carrying opposite judgment outcomes, $n = 23$).

C.1 Who writes the decision

The refusal write is distributed, not a sparse safety-head circuit. Cumulative channel-matched specificity reaches only 45% at the top ten heads and needs 67 heads to reach 80%; the write is led by one head with a long tail, and multilayer perceptrons carry 38% of the decision-site write (write fraction 0.384, below the 0.50 Jacobian threshold so the head decomposition is adequate).

Table 4: Per-head write onto the refusal direction and channel-matched specificity, top ten writers. The lead head L16 H23 alone carries 11.7% of the total specificity; writers span layers 11 to 16, and L15 H15 is the sole anti-refusal writer. Refusal is written broadly into the decision channel, led by one head but not carried by it.

Head	Write onto refusal	Channel-matched specificity
L16 H23	+0.742	+0.756
L15 H2	+0.302	+0.368
L14 H19	+0.334	+0.347
L15 H0	+0.265	+0.285
L11 H20	+0.246	+0.274
L16 H21	+0.172	+0.197
L14 H22	+0.178	+0.193
L15 H6	+0.175	+0.189
L13 H29	+0.139	+0.144
L15 H15	−0.130	−0.142

C.2 What each writer reads

Every one of the ten top writers reads content only weakly aligned with the moral subspace: none clears the moral-family band, and none is a clean harm-copy head. Moral-subspace fraction runs 0.15 to 0.28 with comparable harm loading, and the writers split into instruction-attenders and content-attenders, but all are labeled neither-moral-nor-harm.

Table 5: What each top writer reads, in the shared residual basis. No single head reads the moral subspace cleanly, consistent with the causal result that the moral subspace carries only a specific minority of the refusal effect.

Head	Attention plurality	Instruction-token frac	Content frac	Moral-subspace frac	Harm cosine	Label
L16 H23	instruction	0.735	0.265	0.278	0.264	neither
L15 H2	content	0.302	0.679	0.219	0.226	neither
L14 H19	instruction	0.776	0.222	0.238	0.247	neither
L15 H0	instruction	0.642	0.312	0.254	0.247	neither
L11 H20	instruction	0.875	0.125	0.195	0.188	neither
L16 H21	instruction	0.740	0.260	0.254	0.229	neither
L14 H22	content	0.438	0.555	0.239	0.260	neither
L15 H6	content	0.040	0.637	0.174	0.032	neither
L13 H29	template	0.001	0.019	0.146	0.031	neither
L15 H15	content	0.096	0.510	0.173	0.086	neither

C.3 The decisive interchange cells

Patching the decision channel with a request-twin’s content and reading the induced change in the refusal and judgment projections. The moral-subspace restriction moves refusal about a third as much as the full patch, almost all of which is the harm slice; a random rank-3 patch moves nothing.

Table 6: The decisive interchange cells. The moral subspace is a specific refusal substrate: restricting to it moves refusal more than a random rank-3 patch does ($\Delta = 0.031$, paired 95% CI [0.020, 0.043], excludes 0). But the harm-restricted patch (−0.0261) nearly equals the full moral-subspace patch (−0.0282), so almost all of the specific effect is harm. The ratio-of-ratios of restricted-to-full transfer is wider than the sweep below (refusal 0.34, judgment 0.52, difference 0.18 with a bootstrap 95% CI of [−0.24, 0.39] that includes 0 at this count), so the sweep, not the single ratio, resolves the shape.

Interchange cell	Effect	Minimum detectable effect
Full → refusal	−0.0833	0.0238
Moral-subspace-restricted → refusal	−0.0282	0.0238

Interchange cell	Effect	Minimum detectable effect
Complement (off-subspace) $\rightarrow$ refusal	-0.0636	0.0238
Harm-rank-1 $\rightarrow$ refusal	-0.0261	0.0238
Random-rank-3 $\rightarrow$ refusal (control)	-0.0005	0.0238
Full $\rightarrow$ judgment	+0.0459	0.0086
Moral-subspace-restricted $\rightarrow$ judgment	+0.0237	0.0086

C.4 The nested rank sweep

Restricting the content patch to nested moral subspaces of rank k (the top- k eigenvectors of the paired moral-neutral content-contrast covariance, $k \leq 16$) and reading the transfer coefficient $R_r(k)$, the fraction of the full interchange effect on readout r reproduced under the restriction. This rank- ≤ 16 basis is not the variance-threshold object Section A.1 rejects (the uncentered effective-dimension-385 span at 0.90 pooled-difference variance, on which refusal, persona, and random directions all project 0.7–0.8 and which “discriminates nothing”): its per-rank subspace purity, the fraction of the mean moral contrast the rank- k span captures (tabulated below), holds at 0.97–0.99, and the random-direction null transfers near zero at every rank. So what the sweep transfers is moral-contrast signal certified by that purity, not the generic high-variance content the rank-385 object cannot separate; judgment climbing while refusal saturates on the same purity-0.97–0.99 basis is a difference in what the two readouts read, not judgment tracking a content axis the gate ignores.

Table 7: The nested rank sweep, replicated and pooled. The original 23-twin run and a 19-twin replication (48 newly authored twins, 19 surviving the same baseline-discrimination screen) were run through one session; the original subset reproduces the run of record ($R_{\text{refusal}}(16)$ 0.27, verdict `harm_saturating`), the new subset alone reads `indeterminate` (same sign of the gap, plateau CI touching 0), and the pooled sweep is primary under the pre-registered sign rule. Judgment transfer climbs to 0.66 while pooled refusal transfer rises to $k = 3$ (0.27) and holds flat at 0.22–0.24, at the harm-rank-1 level (harm-rank-1 transfer 0.33), with the random-direction null near zero at every rank. About 76% of refusal’s causal twin-difference input lies outside the rank-16 moral basis (73% already at the rank-3 peak). No per-rank confidence interval on $R_{\text{judgment}}(k) - R_{\text{refusal}}(k)$ is reported; the one difference-CI computed for this contrast (Section C.3) includes 0 at $n = 42$ (0.21, $[-0.07, 0.39]$), as the pre-registered power table predicted (about 140 twins would resolve it).

k	$R_{\text{refusal}}(k)$ pooled, $n = 42$ [95% CI]	original 23	new 19	$R_{\text{judgment}}(k)$	Random- direction null	Subspace purity
1	0.03 $[-0.01, 0.07]$	0.01	0.05	0.05	≈ 0	0.97
3	0.27 $[0.14, 0.48]$	0.31	0.22	0.46	≈ 0	0.98
8	0.22 $[0.11, 0.40]$	0.26	0.16	0.59	≈ 0	0.99
16	0.24 $[0.13, 0.41]$	0.27	0.20	0.66	≈ 0	0.99

C.5 The one-knob fit

The whole sweep collapses to a single free parameter: refusal transfer is judgment transfer clipped at a harm ceiling, $R_{\text{refusal}}(k) \approx \min(\text{harm ceiling}, R_{\text{judgment}}(k))$ with the ceiling 0.25 on the pooled twins (0.28 original, 0.19 new; grid least squares over the four ranks).

Table 8: The one-knob fit on the pooled sweep (RMSE 0.023 over all four ranks, 0.022 on the plateau; on the original 23 twins the ceiling is 0.28 at RMSE 0.027), while two harm-amplitude alternatives (the ceiling scaled by the harm-capture fraction, and by its square) miss by 0.10 to 0.24. The one place it strains is rank 1, where it over-predicts (measured 0.027 against predicted 0.053): the highest-variance contrast component, the most harm-aligned single direction (variance purity 0.974, cosine 0.35 to harm), is causally inert, moving neither readout at rank 1. Variance is not causal relevance; the harm read is a rank-1 causal object that is not the rank-1 variance object.

k	Measured $R_{\text{refusal}}(k)$, pooled	One-knob prediction	Residual
1	0.027	0.053	−0.026
3	0.275	0.246	+0.029
8	0.222	0.246	−0.024
16	0.242	0.246	−0.004

C.6 The harm-partialed identification cell

Harm alone nearly reproduces the full moral-subspace effect, but a small non-harm moral read remains at the edge of the instrument. Projecting the harm direction out of the moral subspace and patching the residual still moves refusal -0.0133 (95% CI $[-0.023, -0.005]$, excludes 0), about half of the full moral-subspace effect. This point estimate is below the refusal interchange minimum detectable effect (0.0238) even though its bootstrap CI excludes 0, so it sits at or near the detection limit and should be read as a boundary result, not a robustly resolved one. The harm direction captures a fraction 0.46 of the moral subspace, and the subspace-versus-complement decomposition is additive (ratio 1.10, 95% CI $[0.91, 1.35]$, includes 1). So the moral subspace’s refusal effect is harm-dominant with a small non-harm residual at the detection limit, and that residual is the one place, on this instrument, where refusal reads moral content beyond harm.

C.7 The behavioral severity ladder

The harm-keyed, saturating read is coherent with OLMo-3 being a weak intent-refuser. On intent-harmful requests its refusal reaches only about 17% at top severity (violating items 0 / 0.17 / 0 / 0.17 / 0.17 across the five-level severity ladder, benign items 0 throughout), so the behavioral operating band is nearly empty. This is a model property (weak coupling between intent severity and refusal) that a harm-surface-keyed gate predicts, not a stimulus artifact, and it is the reason the cross-model commitment axis in Section D is measured on Llama and GPT-OSS rather than on OLMo, whose refusal barely fires on these requests both in projection (-0.08) and in behavior (17%).

C.8 The reconciled cross-ablation (activation-level projection-out)

A second causal instrument, distinct from interchange and from the weight-folded ablation of Section F: the direction is projected out of the layer-16 residual output at every position during generation, and the outcomes are read behaviorally with the chat-template harness of the cross-model battery (100 held-out harmful requests for refusal, 120 forced-choice items for moral judgment; baseline refusal 0.62, judgment accuracy 0.74). The directions are the refusal direction (400/400 train set), the judgment-decision direction, the persona direction, and five random unit directions, all at the same site.

Table 9: The reconciled cross-ablation on OLMo-3-Instruct. Paired bootstrap over requests (2000 draws): ablating the judgment-decision direction raises refusal by 0.12 over the random draws (95% CI [0.06, 0.19]); ablating the refusal direction leaves judgment accuracy within 0.007 of the random draws ([−0.018, 0.033]; no cross-effect detectable at about 0.13 in judgment accuracy). The pre-registered minimum detectable effect for a refusal-rate difference at this count was 0.14, so the 0.12 is detected by its interval and lies below the planning bar. The persona control moves one request in a hundred against a random reference that moves none.

Ablated direction	Refusal rate	Judgment accuracy	Requests re-decided (comply→refuse / refuse→comply)
none (baseline)	0.62	0.742	
judgment-decision	0.74	0.758	12 / 0
refusal	0.78	0.750	33 / 17
persona	0.63	0.725	1 / 0
random (five draws)	0.62 each	0.733 to 0.750	0 / 0 each

Two things about this instrument are stated rather than assumed away. First, projecting the refusal direction out of the residual at one layer does not remove refusal on this model: it re-decides about half the requests in both directions, and every re-decided output is coherent (all 33 new refusals are well-formed opening refusals; all 17 new compliances are “Certainly, however” redirects; length and repetition match baseline). A single-layer projection-out and a weight-folded orthogonalization are different interventions on the same direction, and the removability claims of Section F rest on the latter. Second, the judgment-decision effect is one-directional (12 requests to refuse, none to comply) and specific relative to five random directions and to the persona direction. Whether it should be read as judgment content feeding the refusal decision or as the removal of one more gate-adjacent direction is not settled by this cell; the remaining discriminator (whether the re-decided requests are the ones nearest the refusal boundary at baseline) is priced in the pre-registration and unrun. The interchange results of Section C.3 remain the primary causal claim; this cell adds an ablation-side arrow, held at this scope.

D Cross-model panel: per-model detail

This appendix gives the per-model support behind the two-axis panel in Section 8: the depth-matched Llama-versus-OLMo battery, the Llama patch-layer sweep and boundary-band cell, the Llama anatomy and robustness anomaly, and the GPT-OSS position gate and reversibility cells. Interchange numbers are transfer coefficients $R_r(k)$ as defined in Section C; the asymmetry statistic is $A = (|\text{engage}| - |\text{disengage}|) / (|\text{engage}| + |\text{disengage}|)$, where engage is the effect of adding harmful content and disengage the effect of removing it.

D.1 The decision-site bottleneck across four architectures

Table 10: The decision site is an 8-to-15 effective-dimensional control-token bottleneck on every architecture tested, including a 20B reasoning mixture-of-experts. All rows are measured on one 240-text sample per model (128 for GPT-OSS) at the primary layer; intervals are subsampling intervals over texts; the shuffle reference is the participation ratio after independent column permutations, which preserves every marginal variance and destroys the correlations. The Llama value of record is 10.3 raw and 14.2 standardized at the same position; the 13.5 quoted in earlier drafts as a second position was the standardized read from the decision-anatomy harness, not a different token. GPT-OSS’s harmony decision token is the one panel position where the held-one-out moral band (0.53) survives the covariance-matched null (q95 0.48), so its position validity is stated for decision-direction reads, and the content survival there is noted rather than assumed away.

Model	Decision-site participation ratio [95% CI]	Standardized	Content positions (last / mean)	Fraction of the column-shuffle reference	Position-valid for content projection?
OLMo-3-7B-Instruct	14.7 [14.3, 16.2]	20.3	62.8 / 40.4	0.066	no (band [0.40, 0.47] below null 0.557)
Qwen2.5-7B-Instruct	8.6 [8.2, 9.4]	13.5	42.4 / 32.7	0.041	no
Llama-3.1-8B-Instruct	10.3 [10.1, 11.1]	14.2	97.3 / 26.7	0.050	no
GPT-OSS-20B	9.4 [9.1, 10.7] (raw); 12.8 standardized	12.8	high-dimensional	0.084	valid for decision reads; moral band above null (0.53 vs 0.48)

D.2 Llama reads broad and commits early: the depth-matched battery

The reads-broad verdict and the asymmetry statistic were both first read at Llama’s read layer 16, which is past the layer where Llama commits. Re-running the full cell battery at Llama’s pre-commitment coherent depth (layer 12), and matching OLMo there, resolves both. The read-layer values are kept only to show the collapse.

Table 11: The depth-matched battery. At matched depth Llama reads the moral subspace broadly (refusal transfer 0.85 essentially equal to judgment transfer 0.79, the gap that stays open on OLMo closes), while OLMo stays harm-keyed (refusal 0.43 below judgment 0.53). The asymmetry difference collapses from +1.03 read at each model’s own layer to +0.26 = (−0.28) − (−0.54) read at matched depth 12, so the read-layer +0.82 on Llama was a post-commitment artifact and the asymmetry is a consequence of early commitment, not a separate property.

Quantity	OLMo-3 at layer 12	Llama-3.1 at layer 12	Llama at read layer 16
R_{refusal} (disengage sweep)	0.43	0.85	(denominator-latched)
R_{judgment}	0.53	0.79	—
harm-rank-1 restriction	harm-keyed (gap open)	0.59 (gap closes)	—
Asymmetry A	−0.54 (CI [−0.81, −0.32])	−0.28 (CI [−0.47, +0.03])	+0.82 (CI [0.19, 0.98])
Read verdict	harm-keyed	broad moral	—

The reads-broad verdict survives a harm-coextensive alternative at rank 1: weighting each moral principal component by its marginal contribution to the engage effect, the request-twin harm direction spans only 3.6% of the engage-driving moral basis, with the engage weight sitting on the second and third components (0.23 each) where the harm direction captures 9.4% and 0.03%. A single harm

cue cannot masquerade as the broad read. The rank-2/4 version was then run on the severity-ladder contrasts (30 pairs) and the boundary twins (36 pairs) at layer 12: engage-weighted capture 0.086 / 0.192 / 0.207 (severity) and 0.049 / 0.127 / 0.142 (boundary) at ranks 1 / 2 / 4, against a control basis built the same way from sentiment contrasts at 0.092 / 0.140 / 0.239, syntax and register at or below 0.14, a random-basis q95 below 0.002, and a self-capture positive control of 0.755 for the moral PCs’ own split half. A rank-4 harm basis captures no more of the engage-driving basis than a rank-4 sentiment basis does. The moral principal components were re-derived in-run (the original session saved none); a per-component parity check against the four saved harm cosines matched on components 1, 2 and 4 and missed on the near-degenerate third component (0.157 against 0.018), while the subspace-level parity (the harm projection onto the rank-4 span, 0.336 against 0.367) held within 0.05, and the cell is read under the subspace rule (pre-registration Amendment 16.3).

D.3 Llama patch-layer sweep and boundary cell

Table 12: Llama’s disengage is coherent below the read layer and incoherent at it, so by the frozen rule the verdict is early commitment: the refusal decision crystallizes before layer 16. OLMo’s disengage is coherent at the read layer 16 (−0.62), so OLMo commits at or after the read layer.

Patch layer	Llama disengage effect	Coherent?
8	−0.12	yes (CI excludes 0)
12	−0.11 (full cell −0.57)	yes (CI excludes 0)
14	−0.20	yes (CI excludes 0)
16 (read layer)	−0.014	no

The boundary-band bidirectional cell (36 micro-graded twins at Llama’s roughly 0.5-refusal severity, all three sub-levels inside the [0.4, 0.7] unsaturated band) shows the directional asymmetry directly: engage (add harmful content) moves refusal +0.142 (95% CI [+0.086, +0.212], sign fraction 0.81, coherent), while disengage (remove harmful content) moves it −0.014 (95% CI [−0.084, +0.052], sign fraction 0.51, incoherent). Llama refuses on intent (baseline refusal 9/10, operating band severity 3 to 5), so its refusal cell is measurable where OLMo’s is empty.

D.4 Llama anatomy and the robustness anomaly

Llama’s anatomy is OLMo-like: pre-norm reconstruction 1.0008 (no fold needed, an architecture cross-check), a clean low-dimensional decision channel (decision-token-harness participation ratio 13.5, with the in-format value of record 10.2 in Section D.1; the covariance null moving only 0.148 to 0.114 under standardization, so the dim-788 outlier lives at content positions and not at the decision bottleneck), a distributed write with a 30% multilayer-perceptron share, and all top writers labeled neither-moral-nor-harm.

Llama is the panel’s robustness anomaly: its refusal is entangled with moral judgment where the other models’ is not. At the best ablation layer, removing refusal drops judgment accuracy from 0.75 to 0.604, far outside the matched-random-ablation band (0.747 ± 0.007) and dose-dependent (Spearman 1.0); refusal removability is also family-dependent, dropping only from 0.900 to 0.475 on Llama against clean removal on OLMo and Qwen. Early commitment of a broad moral read is the mechanism: because Llama reads broad moral content and commits before the decision site, ablating its refusal reaches into the moral read in a way OLMo’s harm-keyed late-committing gate does not.

D.5 GPT-OSS position gate and reversibility

GPT-OSS contributes three cells that do not depend on the causal interchange (held for this model), plus the reversibility result.

Table 13: GPT-OSS Tier-1 cells. The position gate is the strongest cross-model generalization of the bottleneck finding (a 20B reasoning mixture-of-experts) and licenses the projection reads. GPT-OSS is a reversible reader: an inculcating prefill flips benign requests to refuse 7/7, and a graded exculpatory prefill flips ceiling-refusing violating items to comply 6/10 with the decision-channel projection moving monotonically toward comply in all 10 items. The behavioral flip is the primary evidence; the projection corroborates with the caveat that it reads the last token of the injected prefill. The reads-harm placement is correlational (the prompt-to-trace harm-loading), not causal, because the interchange sweep is held for this model. The graded structure is the control that rules out the prefill merely asserting benignness: the weakest prefill does not read lowest, the projection spikes (215) then falls (78) as rhetorical strength climbs while the behavioral flip rate rises. The within-harm-status commitment curve is not computable at this operating point (the gate is a step function, only 5.6% of violating items in the mid-band), and is reported as such rather than replaced by a harm-separability fallback.

Cell	Value
Position gate (harmony decision channel)	participation ratio 12.8, below the 25 ceiling, position-valid
Engage flip (inculcating prefill, benign → refuse)	7/7 (Wilson 95% [0.65, 1.0])
Disengage flip (graded exculpatory prefill, violating → comply)	6/10
Decision-channel projection under graded disengage	moved toward comply monotonically (10/10 at the prefill token, 8/8 at the post-response decision token), about −1 SD of the sample; not distinguishable from covariance-matched random directions (one-sided p 0.17 to 0.23), so not a corroborating leg
Decision-channel null-ratio	372 (the channel’s dominant axis of variation is the refusal split)
Prompt harm-loading (instruction token)	cosine 0.977 to harm against 0.001 harm-orthogonal (near-purely harm)
In-trace harm-loading	cosine 0.49 to harm against 0.13 harm-orthogonal (harm-dominant, attenuated)

D.6 The two-axis mapping

The two axes are *what* refusal reads (harm versus broad moral content) and *how* it commits (at the read layer, early, or reversibly). OLMo reads harm by interchange (transfer saturates at the harm-rank-1 level, ceiling 0.31) and commits at or after the read layer (disengage coherent there, −0.62). Llama reads broad moral content by interchange at matched depth (refusal transfer 0.85 essentially equal to judgment 0.79) and commits early (disengage coherent below layer 15, incoherent at the read layer 16). GPT-OSS reads harm correlationally (prompt cosine 0.977 to harm against 0.001 orthogonal, causal test held) and is a reversible reader on behavior (engage 7/7, disengage 6/10, 5/10 on replication). Qwen reads beyond the harm-rank-1 level by interchange (refusal transfer 0.54 at rank 16 against harm 0.38, gap to judgment 0.12 unresolved at 19 twins, verdict indeterminate) and commits bidirectionally at the read layer (disengage −2.70, engage +0.68, both coherent). The table is the measured result; its interpretation as a dimensionality-to-reversibility law is a hypothesis on three architecture-confounded points plus one indeterminate point, stated for testing in Section 10, not a mechanism established.

E Reproducibility

E.1 Hardware

The causal interchange and per-head attribution on the 7B instruct models, and the GPT-OSS-20B reasoning cells, run on a single NVIDIA A100-80GB. The zero-GPU analysis (bootstrap confidence intervals, the rank sweep and one-knob fit from saved arrays, the calibration ladder, and figure regeneration) runs on CPU, or on Apple Silicon via the PyTorch MPS backend where a GPU is not required. Once the per-head, per-request-twin, per-rank, and per-rollout arrays listed in Section E.7 are saved, every statistic in the paper is re-derivable without GPU access.

E.2 Random seeds

Purpose	Seed	Where set
Moral probing train/test split	42	<code>build_probing_dataset(dataset_version="v2")</code>
Bootstrap resampling (bands, interchange cells, reversibility)	0	per-cell statistics harness, $B = 2000$ resamples
Covariance-matched null directions	fixed per subspace	deterministic from the realized subspace geometry (Section B.2)

All bootstrap confidence intervals use $B = 2000$ resamples at seed 0; the null and persona controls are deterministic functions of the constructed subspace and are realized before any refusal vector is projected.

E.3 Software

- Python 3.13
- PyTorch (with CUDA on the A100 runs, MPS on Apple Silicon for zero-GPU analysis)
- HuggingFace transformers 5.12.1
- HuggingFace datasets
- numpy, matplotlib (figure regeneration)

Exact versions are pinned in `pyproject.toml` in the released codebase; the standard repository environment is used without experiment-specific overrides. GPT-OSS-20B is loaded from its mxfp4 weights and dequantized to bf16 for probing.

E.4 Model checkpoints

The four panel models and their extraction geometry:

Model	Repo	Layers	Read layer	Used for
OLMo-3-7B base	allenai/olmo-3-1025-7B	32	16	crystallization trajectory, base proto-refusal
OLMo-3-7B instruct	allenai/olmo-3-7B-Instruct	32	16	causal interchange, per-head attribution, bottleneck, orthogonality
Qwen2.5-7B instruct	Qwen/Qwen2.5-7B-Instruct	28	14	bottleneck, decision orthogonality
Llama-3.1-8B instruct	meta-llama/Llama-3.1-8B-Instruct (gated)	32	16	bottleneck, broad-read + early-commitment panel
GPT-OSS-20B (reasoning MoE)	openai/gpt-oss-20b	24	12	position gate, reversibility, in-trace harm-loading

Additional models appear in the cross-model harm-direction cells (Section A.4): allenai/Olmo-3-7B-Think (reasoning variant, 32 layers, in-trace refusal gradient), deepseek-ai/DeepSeek-R1-Distill-Llama-8B and deepseek-ai/DeepSeek-R1-Distill-Qwen-14B (the reasoning distills, in-subspace fraction 0.11), and Qwen/Qwen2.5-14B-Instruct and meta-llama/Llama-3.1-8B-Instruct (reply-inversion causal validation). GPT-OSS-20B is a mixture-of-experts (32 experts, 4 active per token). Checkpoints load from each repo’s default branch at the pinned transformers version; drivers read the live layer count and hidden size and fail loud if a re-release changes the geometry, so a silently re-pinned repo cannot mis-index a band.

E.5 Figure regeneration

The three flagship figures produced for this paper (the crystallization pairing, the OLMo one-knob sweep, and the GPT-OSS reversibility panel) regenerate from committed CSVs with no model, GPU, or network access:

```
# From papers/fl_what_refusal_reads/figure_data/
python3 regen_fl_figures.py
```

Each figure reads only its committed CSV in figure_data/ and writes a vector PDF to figures/: fl_crystallization.csv to fl_crystallization.pdf, fl_one_knob.csv to fl_one_knob.pdf, and fl_gpt_oss_reversibility.csv to fl_gpt_oss_reversibility.pdf. The CSVs carry the verified committed values (for example the sweep CSV holds R_{judgment} 0.05/0.46/0.59/0.66 and R_{refusal} 0.01/0.31/0.26/0.27 against a random null of 0 and the one-knob fit 0.05/0.31/0.31/0.31). The three remaining figures (the bottleneck participation-ratio bars, the calibration ladder, and the depth-collapse panel) are committed vector PDFs regenerated by the same convention. Colors follow a fixed-semantic Material palette (moral and judgment indigo, refusal red, valid and comply green, null and reference gray), assigned consistently across every figure and never carried by color alone, so identity survives grayscale printing and color-vision deficiency. Figures are byte-reproducible: the regeneration scripts pin SOURCE_DATE_EPOCH, so re-running them produces identical PDFs.

E.6 Pre-registration and amendment trail

The credibility of the null and threshold verdicts rests on their being frozen before the numbers were computed. Each of the three study directions behind this paper carries a pre-registration document committed to the public repository before its headline quantities were run, with a dated amendment trail: the moral-subspace orthogonality rule (the decision margin $M = 0.05$, the conjunction of null and control bars, and the naming of the null outcome as the more publishable one) was fixed before any moral subspace or refusal vector existed; the calibration pre-registration (the ladder, the band construction, and the paired band-minus-projection test) was committed before any calibrated headline was computed; and the causal anatomy pre-registration fixed the interchange transfer coefficients, the minimum-detectable-effect reporting rule, and the both-branches framing of every shape verdict before the compute session. Post-hoc analysis-choice changes are recorded as dated amendments committed before the recompute they license, so a reader can separate what was predicted from what was discovered.

E.7 Artifact save-list and public release

Per the project’s reproducibility convention, per-unit arrays are saved so that all statistics stay re-derivable without re-running a model:

- **Per-head** write-onto-refusal and channel-matched-specificity arrays, and per-head read vectors in the shared residual basis (the tables in Section C.1, Section C.2).
- **Per-request-twin** paired interchange deltas for the refusal and judgment readouts, for the full, moral-subspace-restricted, complement, harm-rank-1, and random-rank-3 patches (the cells and the paired specificity confidence interval in Section C.3).
- **Per-rank** transfer coefficients $R_{\text{refusal}}(k)$, $R_{\text{judgment}}(k)$, and the random-direction null across $k \in \{1, 3, 8, 16\}$, with per-rank purity (the sweep and one-knob fit in Section C.4, Section C.5).

- **Per-rollout** GPT-OSS decision-channel refusal projections across the graded-prefill series (the graded-projection cell in Section D.5, including the post-response decision-token reads).
- **Per-pair** moral-neutral content-contrast difference vectors and the covariance-matched null resample arrays behind the calibration ladder (Section B.1).

Output JSON carries full metadata per record (model name and repo, extraction position and layer, participation ratio, format, number of pairs, harness and classifier version, and timestamp), so a result traces to the artifact it came from. All code, scripts, committed figure CSVs, and structured output are released at <https://github.com/deepsteer/deepsteer/>.

The distilled artifacts behind every figure and headline number, the per-head write and specificity arrays, the calibrated covariance nulls, the participation-ratio profiles, the calibration ladders, and the rank-sweep outcomes, are indexed in `deepsteer/supplement/MANIFEST.json`, each with a content hash, its provenance, and the figure or table it backs. The two instruments shared with the companion methods note (the decision-site participation-ratio profile and the depth-asymmetry panel) live in the supplement once and are cited by both papers; `deepsteer/supplement/scripts/verify.py` checks that this paper’s plotting copies carry the same values as the canonical files, so a shared number can change in only one place. Model ids, decision layers, standardization settings, and seeds are pinned in `deepsteer/supplement/PROVENANCE.md`. The per-unit arrays behind every confidence interval in this paper (per-twin interchange deltas for every cell and rank, per-head write and specificity arrays, per-prompt ablation outcomes, per-rollout position activations, per-pair contrast diffs and null resamples, the decision-site activation samples, and the per-checkpoint proto-refusal directions) are deposited on Zenodo as `deepsteer_fl_arrays_v1.tar.zst` (336 files, 2.8 GB uncompressed) under CC BY 4.0, with a manifest that records each file’s SHA-256, its unit and model, the Hugging Face commit hash of the model that produced it, and the figure or claim it backs (DOI: 10.5281/zenodo.22731361; built by `deepsteer/supplement/scripts/build_release.py` at the deposit commit). Caches whose stimuli are MORABLES retellings inherit that corpus’s non-commercial license and are excluded from the public record with a regeneration recipe; the headline arrays use Moral Stories, public-domain fable retellings, and ETHICS, and are all deposited. Model weights are never deposited.

F Removability battery

The behavioral and interventional companion to the representational results folds into this and the next two appendices. This appendix carries the single-direction refusal-ablation battery across the three instruct families, the behavioral side of the representational dissociation in Section 4, and the mechanism-level reading of Section D.4. Every number here is a measured quantity from the program’s runs; the appendix reports the battery and states what it does and does not license.

F.1 The battery and its readouts

Single-direction refusal ablation follows the Arditi construction: the refusal direction is the difference of mean residual activations between harmful and harmless requests at the last prompt token, and it is orthogonalized out of the attention output projection and the multilayer-perceptron down projection at one layer, chosen per model by a depth-fraction sweep (OLMo layer 19, depth 0.59; Qwen layer 14, depth 0.50; Llama layer 13, depth 0.41). Four readouts are taken before and after: the refusal rate on harmful requests, fresh per-foundation probe accuracy for moral content, the effective dimensionality of the moral subspace, and behavioral moral judgment on the 48-scenario forced-choice battery. Persona-shift compliance, a fifth readout, is reported in Section H.

The representational readouts do not move. On all three models the ablation leaves probe accuracy at 1.0 and the effective dimension of the moral subspace at 5, before and after. On OLMo-3 the ablated direction itself projects only 0.10 of its norm into the moral subspace (mean absolute cosine 0.06 to the foundation directions), and removing it takes refusal from 0.25 to 0.00 on the Paper-5 borderline-plus-harmful request set while leaving comprehension where it was (base-to-fresh cosine 0.749, probe accuracy 1.0, effective dimension 5) and behavioral moral judgment essentially unchanged (0.73 against 0.75). The two OLMo refusal baselines quoted in this paper come from two request sets by construction: 0.25 on the Paper-5 borderline-plus-harmful set, and 0.575 on the Paper-6 held-out harmful set used for the cross-model battery below; each ablation result is read against its own baseline.

Removability is family-dependent. On the held-out harmful set, the single-direction ablation takes OLMo from 0.575 to 0.000 and Qwen from 1.000 to 0.000, but Llama only from 0.900 to 0.475. A rank-one edit removes refusal cleanly on two families and only halves it on the third.

F.2 Behavioral judgment under ablation

Behavioral moral judgment is the readout that separates the families a second time. Under the same ablation, OLMo’s judgment accuracy moves from 0.75 to 0.79 and Qwen’s from 0.875 to 0.812, both within the range of the matched-random control; Llama’s drops from 0.75 to 0.604. The Llama drop sits far outside the magnitude-matched random-ablation band (0.747 with a spread of 0.007 across random directions of the same norm), and it is dose-dependent: with the ablation scaled to half strength the judgment drop is already 0.083 (95 percent bootstrap interval [0.02, 0.17], 48 items) while the refusal rate is still 0.90, and at full strength the drop is 0.146 ([0.06, 0.25]). The matched-random null stays at 0.74 even at twice the full-ablation magnitude, and ablating the persona direction (Section H) leaves judgment at 0.75. So the Llama judgment drop is specific to the refusal direction, not to any salient direction of that norm, and it appears before the refusal rate has moved.

The anomaly resolves upstream rather than in the battery. The depth-matched Llama cell battery of Section D.2 finds that Llama’s refusal reads broad moral content by interchange (refusal transfer 0.85 against judgment 0.79 at layer 12) and commits early. A refusal direction estimated on a broad-moral reader carries moral content with it, so removing it costs judgment; on the harm-keyed readers it does not. The battery is the behavioral face of the reads axis in Section 8.

F.3 What removability does and does not show

On OLMo the ablated direction reads a rank-1 harm slice: 76 percent of refusal’s causal input lies outside the rank-16 moral basis (Section 7), which is why a rank-one edit removes it without touching comprehension. That is a causal statement about one model. Across the panel the link between a low-rank read and clean removability is correlational: two harm-keyed readers are cleanly removable and one broad reader is not, and the read axis on the third harm-keyed model (GPT-OSS) is itself correlational (Section G). This appendix reports the battery; it does not claim the low-rank-read-to-removable link is general, and the Limitations section (Section 10) carries that scope alongside the architecture confound on the two-axis table.

A second scope note comes from the reconciled cross-ablation run on OLMo-3-Instruct (Section C.8): a single-layer projection-out of the refusal direction, applied at the layer output for every position rather than folded into the weights, does not remove refusal on that model. It re-decides about half of 100 held-out harmful requests in both directions (33 from comply to refuse, 17 from refuse to comply) with fully coherent text, while five random directions of the same norm at the same site re-decide none. The weight-folded Arditi ablation of this appendix and the activation-level projection-out are different instruments with different outcomes on the same direction, and the paper keeps them apart: removability claims rest on the former.

G Distributed refusal

Two independent observations say that refusal is a distributed write rather than a single-direction switch: the GPT-OSS held-out ablation battery and the OLMo-3 per-head write attribution into the decision channel (the anatomy behind Section 7 and Section C). They are measured with different instruments on different architectures and reach the same structural reading.

G.1 GPT-OSS: no single direction ablates refusal

On GPT-OSS-20B no single direction ablates refusal on a held-out, category-diverse request set. The end-of-prompt refusal direction, estimated on a category-spanning training draw, coherently flips 4 percent of held-out refusals; the direction estimated at the last chain-of-thought token flips none; and the direction estimated as the chain-of-thought mean removes refusal in 88 percent of cases only by driving generation into incoherence, which the coherence filter excludes. A representation that no single direction removes is not a bottleneck for refusal on this model.

That has a consequence for the causal program. The single-subspace load-bearing test used on OLMo and Llama (patch the moral subspace, read the refusal change) presumes a channel that a low-rank edit can move. On GPT-OSS that presumption fails at the ablation stage, so the GPT-OSS read axis in Section 8.2 is correlational (a projection read at the instruction token and in-trace), and the causal interchange version stays held. The same fact appears from the other side in Section D: GPT-OSS’s harmony decision token is the one panel position where the held-one-out moral band survives its covariance-matched null (0.53 against 0.48), so content survives there in a way the chat decision sites do not show.

G.2 OLMo-3: a distributed write into a narrow channel

On OLMo-3-Instruct the refusal decision is written into the roughly 13-dimensional decision-site channel by a distributed set of attention heads. Ranked by channel-matched specificity, the lead head (layer 16, head 23) carries 11.7 percent of the total; cumulative specificity reaches 45 percent at the top ten heads and needs 67 heads to reach 80 percent (the saved sparsity curve of record; the per-head arrays are in the supplement as `head_attribution.csv`). Writers span layers 11 to 16. The top ten by specificity are L16 H23 (write +0.742, specificity +0.756), L15 H2 (+0.302, +0.368), L14 H19 (+0.334, +0.347), L15 H0 (+0.265, +0.285), L11 H20 (+0.246, +0.274), L16 H21 (+0.172, +0.197), L14 H22 (+0.178, +0.193), L15 H6 (+0.175, +0.189), L13 H29 (+0.139, +0.144), and L15 H15 (−0.130, −0.142), the sole anti-refusal writer in the top ten.

Attention is not the whole write. Multilayer perceptrons contribute 38 percent of the decision-site write (fraction 0.384), below the 0.50 threshold at which a Jacobian stage would be required and above the 0.23 the un-folded run had reported. That un-folded number is a calibration lesson in its own right: OLMo-3’s reordered normalization makes the naive per-head attribution overshoot, and folding the per-layer RMSNorm gain brings the Stage-1 reconstruction from 3.05 to 0.9999 within a two-sided band of [0.90, 1.10], exact to one part in a billion (methods note, A3).

None of the ten top writers reads moral content in a way the calibrated instruments recognize. All ten are labeled neither-moral-nor-harm: none clears the moral-family band, none is a clean copy-head for harm, and their moral-subspace fractions sit at 0.15 to 0.28 with comparable harm loading. Llama-3.1’s anatomy is OLMo-like on every count that can be compared (pre-norm reconstruction 1.0008 with no fold needed, a distributed write, multilayer-perceptron share 0.30, all writers neither), as Section D.4 reports.

G.3 Reading the two together

On both models the refusal write is many small contributions into a low-rank control channel, a refusal cone rather than a single direction, consistent with the 9-to-15-dimensional bottleneck of Section 5. The GPT-OSS battery shows the consequence for removability (no single direction suffices on a model whose refusal is spread across the trace), and the OLMo attribution shows the mechanism on a model where a single direction does suffice: even there, the direction that a rank-one edit removes is assembled by dozens of heads, none of which carries it alone. The distributed write is what any intervention that widens the read (Section 9) would have to move.

H Persona, the assistant axis, and persona-shift compliance

The persona direction is the program’s named reference axis, a moral-adjacent voice reference rather than a non-moral control; persona-shift compliance is the behavioral battery that reads refusal removal from the compliance side. This appendix collects both so that the calibration ladder (Section B) and the removability battery (Section F) can cite one place. The assistant-axis literature is the framing for the persona direction; no number from it enters this paper.

H.1 The persona direction: decodable, moral-adjacent, not moral

The persona direction is the difference of means between texts written in the assistant’s voice and matched texts in a neutral voice, extracted with the same pipeline as the moral directions. A linear persona probe is highly decodable at every OLMo-3 training stage (peak accuracy about 0.94), while the direction stays nearly orthogonal to the moral foundations: its mean absolute cosine to the

foundation directions rises only from 0.076 at the base model to 0.085 at the Instruct model. Persona is present and stable, and it is not moral content.

It is moral-adjacent. On the rank-3 moral subspace the persona reference projects 0.51 on both the base and the instruct model, just below the moral-family band, which is why it is named a moral-adjacent voice reference in the ladder rather than a non-moral control. The methods note’s calibration case study uses exactly this fact: a reference that projects 0.51 is a rung, not a floor. On GPT-OSS the moral-to-persona cosine is higher (0.30 against OLMo’s 0.24), a general entanglement on that model that raises its persona rung to 0.60.

H.2 Persona-shift compliance under refusal ablation

Persona-shift compliance measures how often the model complies with borderline requests when the request is framed under different persona instructions; the gap between framings is the persona-shift gap. Under single-direction refusal ablation the compliance rate rises on every model, OLMo from 0.75 to 1.00, Qwen from 0.90 to 1.00, and Llama from 0.70 to 0.95, and on OLMo every persona gap closes toward zero. The construction of the battery (borderline requests under four persona framings) follows the cross-model paper’s appendix; the per-cell counts are those of that battery and are not restated here.

Before any intervention, comprehension and compliance are only weakly coupled on OLMo-3: the probability of complying given that the model comprehends the moral content is 0.77, against 0.73 given that it does not. Along the alignment trajectory the agreement between the internal foundation reading and behavioral judgment rises from 0.375 to 0.479 to 0.500 across SFT, DPO, and Instruct, with the corresponding phi coefficient moving from -0.19 to $+0.02$ to $+0.05$. Alignment increases the agreement a little; it does not make behavior a function of comprehension.

H.3 What the persona axis is for in this paper

Three roles. First, a reference rung on the ladder (Figure 2): refusal sits below persona on every model, including the GPT-OSS in-trace peak (0.52 below the persona rung of 0.60), so even the program’s highest refusal projection is less moral-adjacent than a voice reference. Second, a named control in the ablation battery: ablating the persona direction leaves Llama’s judgment at 0.75, so the Llama judgment drop under refusal ablation (Section F) is refusal-specific and not a property of any salient direction. Third, a named control in the reconciled cross-ablation on OLMo-3 (Section C.8), where ablating persona re-decides one of 100 held-out requests against zero for random directions and twelve for the judgment-decision direction.

Scope: the ladder still lacks a non-moral positive-projection control, since persona is moral-adjacent by construction. That is stated as a limitation (Section 10) and is not closed by this appendix.