


Reasoning about the stakes before acting lowers the violating choice more than a same-length non-moral task does, on both models that carry the gap (separate scales; sizes not compared)

(a) OLMo-3-7B-Instruct: truncated reasoning (8% finish)



(b) Llama-3.1-8B-Instruct: mostly completed (90% finish)

