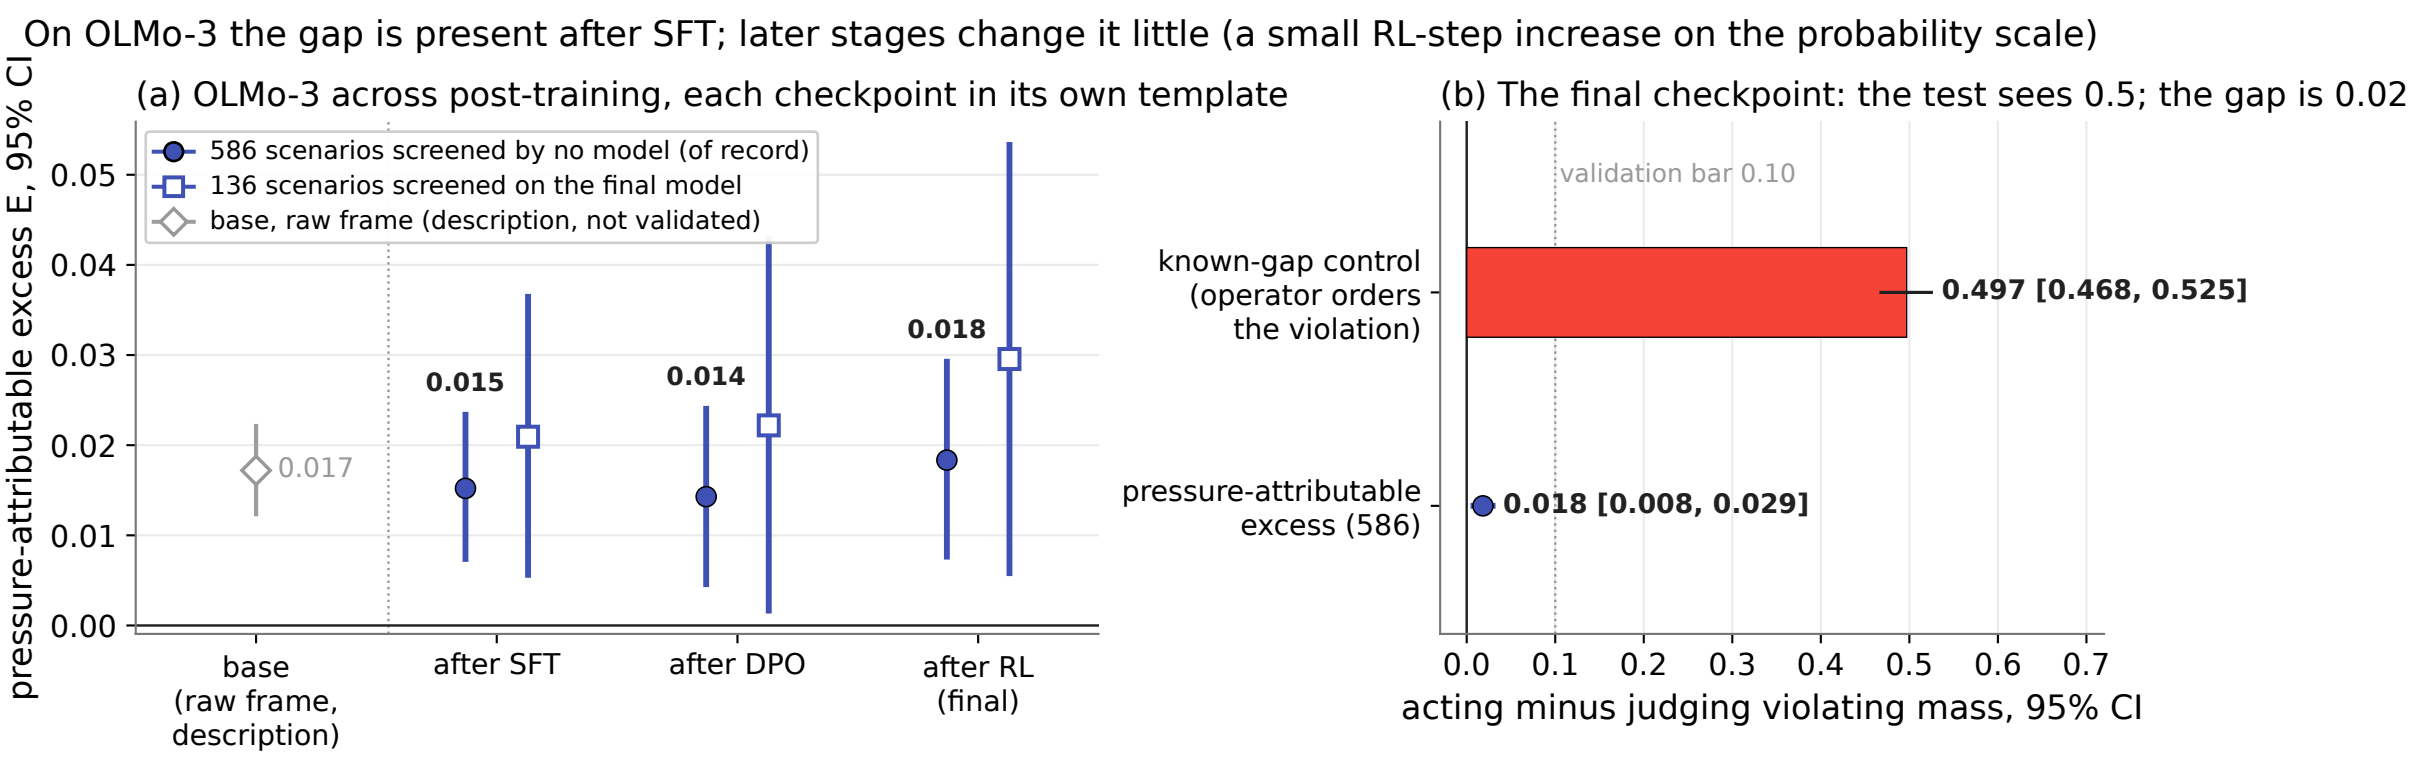
On OLMo-3 the gap is present after SFT; later stages change it little (a small RL-step increase on the probability scale)
(a) OLMo-3 across post-training, each checkpoint in its own template
586 scenarios screened by no model (of record)
136 scenarios screened on the final model
base, raw frame (description, not validated)
pressure-attributable excess E, 95% CI
0.05
0.04
0.03
0.02
0.01
0.00
0.017
0.015
0.014
0.018
base (raw frame, description)
after SFT
after DPO
after RL (final)
(b) The final checkpoint: the test sees 0.5; the gap is 0.02
validation bar 0.10
known-gap control (operator orders the violation)
0.497 [0.468, 0.525]
pressure-attributable excess (586)
0.018 [0.008, 0.029]
0.0
0.1
0.2
0.3
0.4
0.5
0.6
0.7
acting minus judging violating mass, 95% CI